

Big Data

(Hadoop)

Instructor: Thanh Binh Nguyen

September 1st, 2019

S³Lab

Smart Software System Laboratory



“Big data is at the foundation of all the megatrends that are happening today, from social to mobile to cloud to gaming.”

– Chris Lynch, Vertica Systems



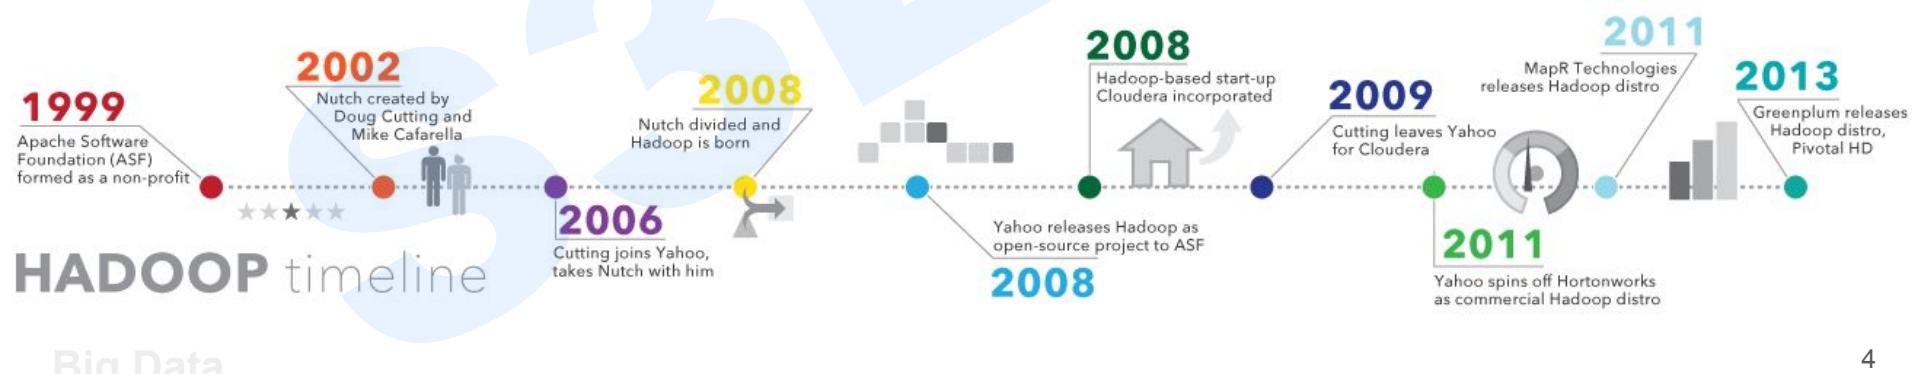
Introduction

- Hadoop is a software framework for distributed processing of **large datasets** (terabytes or petabytes of data) across **large clusters** (thousands of nodes) of computers. Included some key components as below:
 - **Hadoop Common**: common utilities
 - **Hadoop Distributed File System (HDFS)** (Storage Component): A distributed file system that provides high-throughput access
 - **Hadoop YARN** (Scheduling): a framework for job scheduling & cluster resource management (available from **Hadoop 2.x**)
 - **Hadoop MapReduce** (Processing): A yarn-based system for parallel processing of large data sets



Introduction

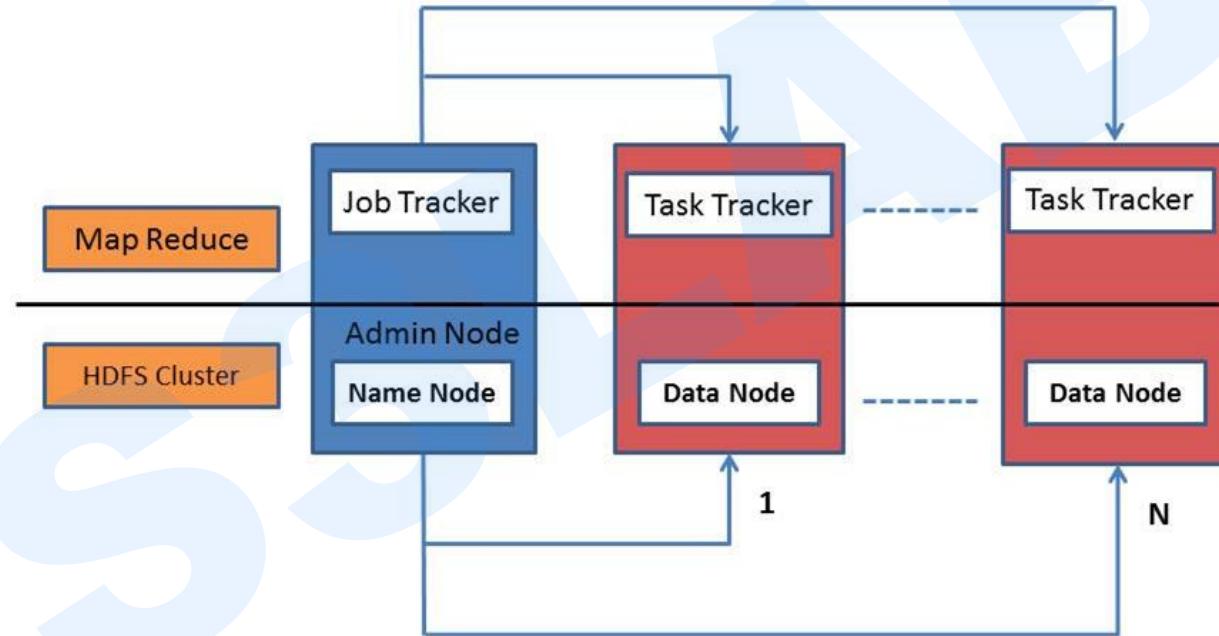
- Hadoop is a large and active ecosystem.
- Hadoop emerged as a solution for big data problems.
- Open source under the friendly Apache License
- Originally built as a Infrastructure for the “Nutch” project.
- Based on Google’s mapreduce and google File System.





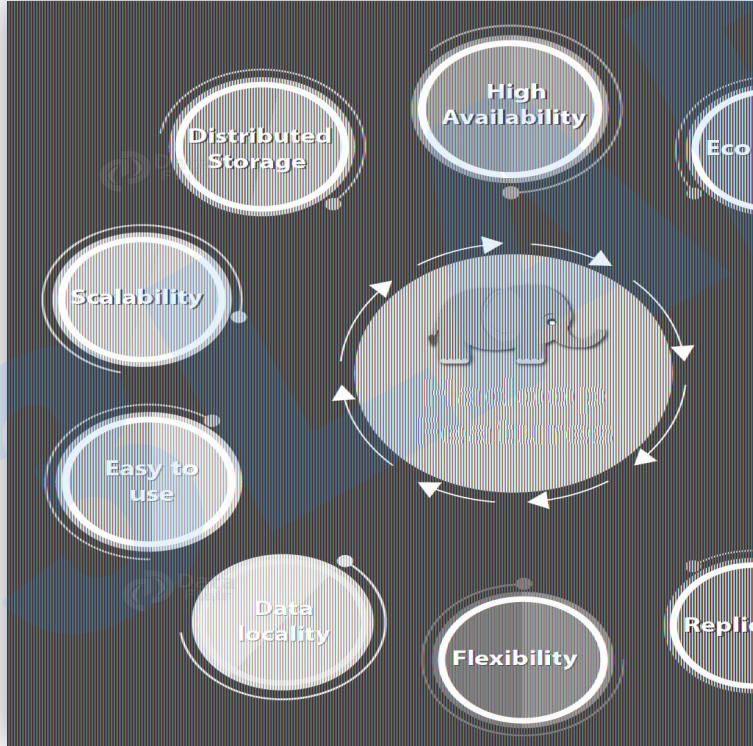
Introduction

Core components





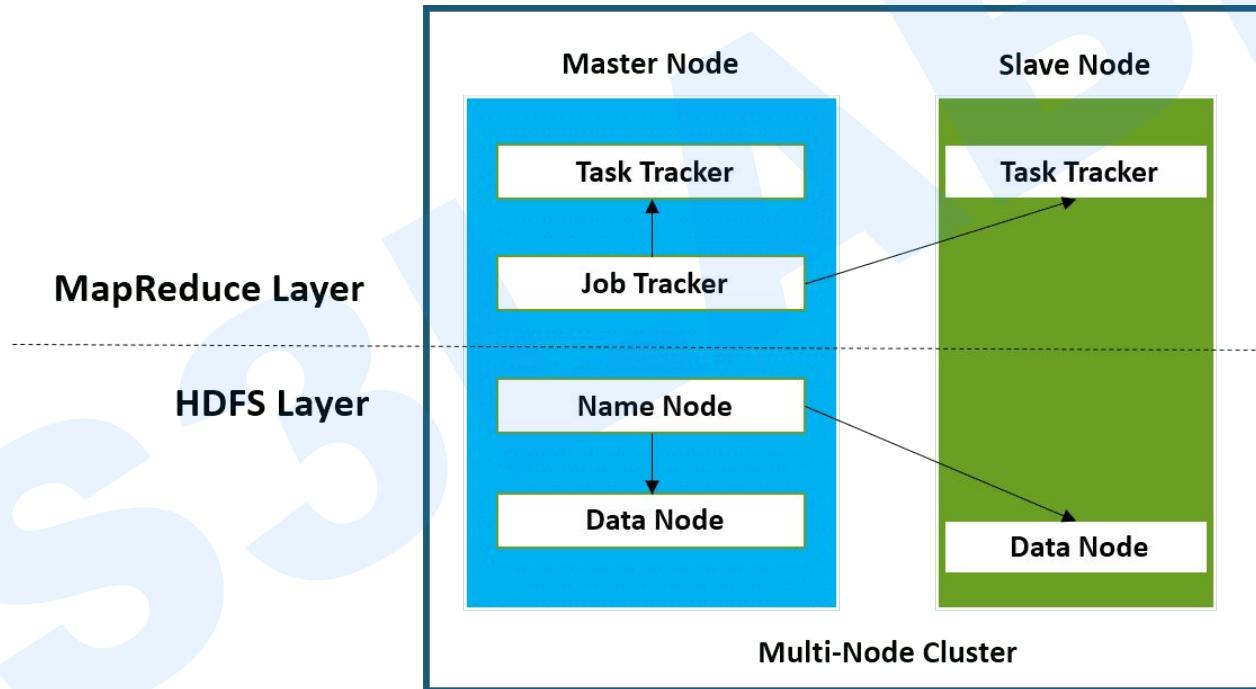
Features





Architecture

Multi-Node Cluster





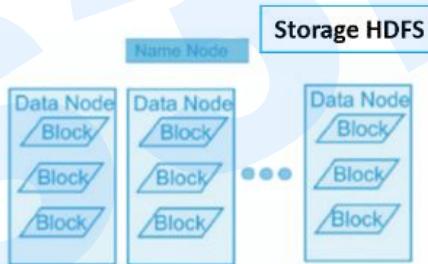
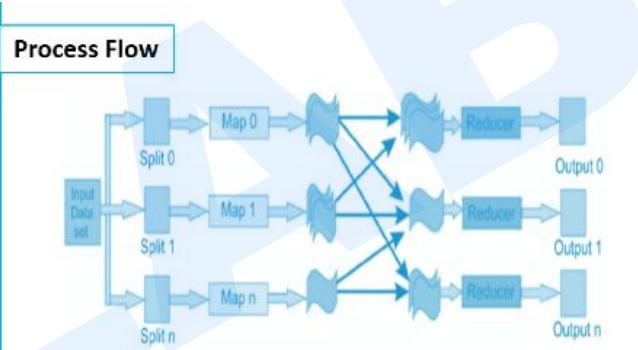
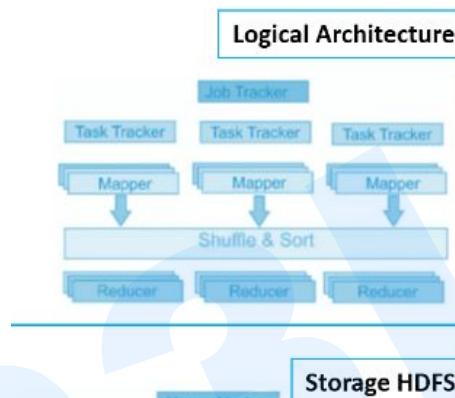
What makes Hadoop unique

- **Data locality and Shared Nothing:** Moving computation to data, instead of moving data to computation. Each node can independently process a much smaller subset of the entire dataset without needing to communicate with one another.
- **Simplified programming model:** allows user to quickly write and test
- **Schema-on-read system** (same as NoSQL platforms) #
Schema-on-write system
- Automatic distribution of data and work across machines

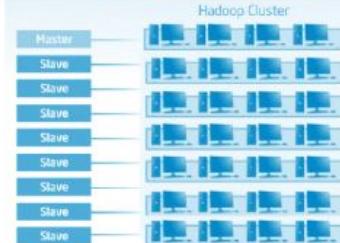


Architecture

Architecture in different perspective



Physical Architecture





HDFS

- **Hadoop Distributed File System (HDFS)** is designed to reliably store very large files across machines in a large cluster. It is inspired by the Google File System.
- Distribute large data file into blocks
- Blocks are managed by different nodes in the cluster
- Each block is replicated on multiple nodes



HDFS

- NameNode:
 - Master of the system, daemon runs on the master machine
 - Maintains, monitoring and manages the blocks which are present on the DataNodes
 - records the metadata of the files like the location of blocks, file size, permission, hierarchy etc.
 - captures all the changes to the metadata like deletion, creation and renaming of the file in edit logs.
 - It regularly receives heartbeat and block reports from the DataNodes.



HDFS

- All of the Hadoop server processes (daemons) serve a web UI. For NameNode, it was on port 50070.

The screenshot shows a web browser window displaying the HDFS NameNode information at localhost:50070/dfshealth.html. The page has a green header bar with tabs for Hadoop, Overview, Datanodes, Snapshot, Startup Progress, and Utilities. The main content area is titled "Overview 'localhost:9000' (active)". It displays the following system information:

Started:	Sun Apr 06 15:52:11 IST 2014
Version:	2.3.0, r1567123
Compiled:	2014-02-11T13:40Z by jenkins from branch-2.3.0
Cluster ID:	CID-5edbd0da-c69f-425b-bbc7-a662ac5d45dc
Block Pool ID:	BP-1127675761-127.0.1.1-1396692597591

Below this is a "Summary" section with the following status information:

- Security is off.
- Safemode is off.
- 35 files and directories, 17 blocks = 52 total filesystem object(s).
- Heap Memory used 34.01 MB of 88.5 MB Heap Memory. Max Heap Memory is 889 MB.
- Non Heap Memory used 40.17 MB of 40.69 MB Committed Non Heap Memory. Max Non Heap Memory is -1 B.

At the bottom, there is a table showing capacity information:

Configured Capacity:	91.54 GB
-----------------------------	----------



HDFS

- DataNode:
 - DataNode runs on the slave machine.
 - It stores the actual business data.
 - It serves the read-write request from the user.
 - DataNode does the ground work of creating, replicating and deleting the blocks on the command of NameNode.
 - After every 3 seconds, by default, it sends heartbeat to NameNode reporting the health of HDFS.



HDFS

- Distributed file system for redundant storage
- Designed to reliably store data on **commodity hardware**
- Built to expect hardware failures
- Intended for
 - Large files
 - Batch inserts

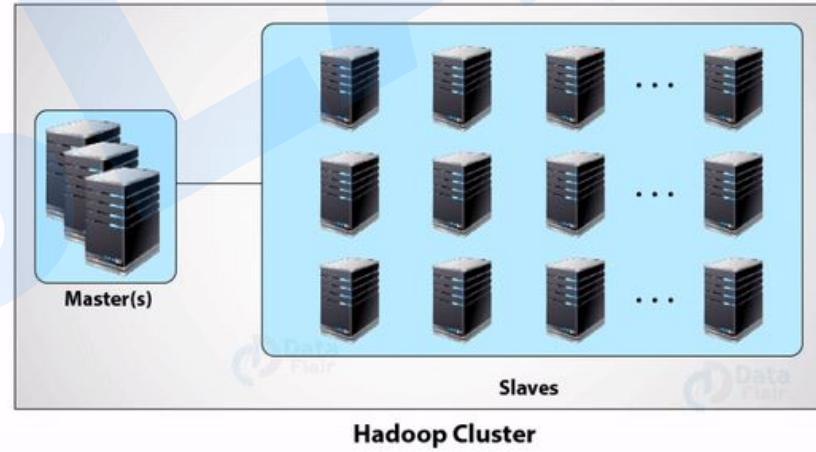


HDFS

Architecture



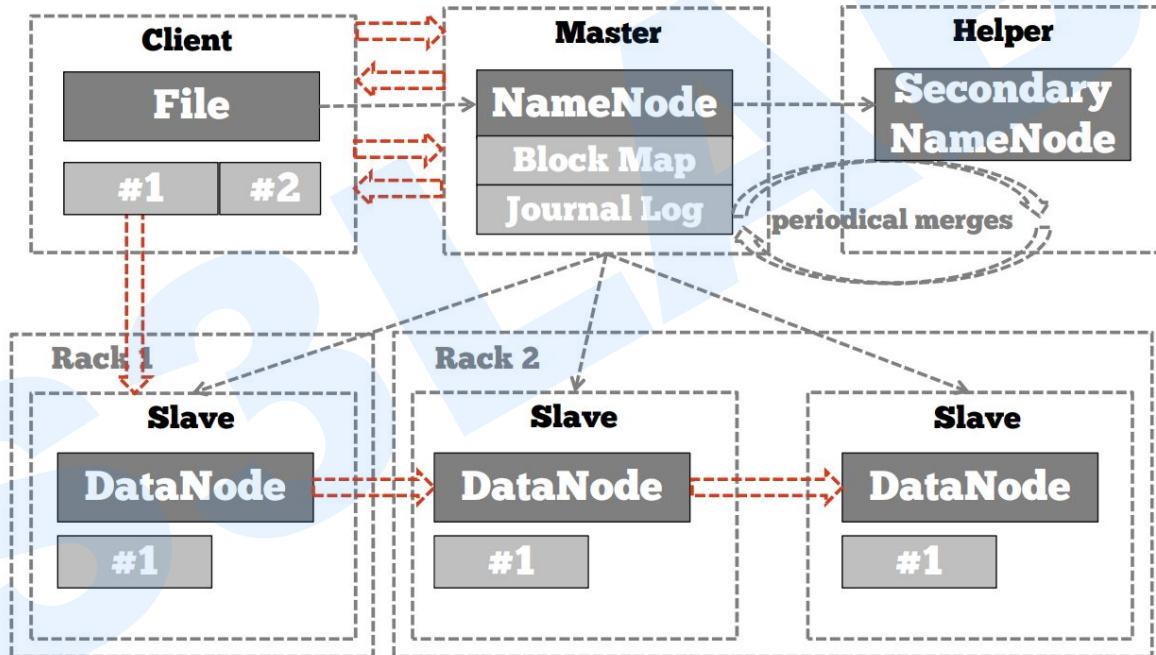
Data Storage in HDFS





HDFS

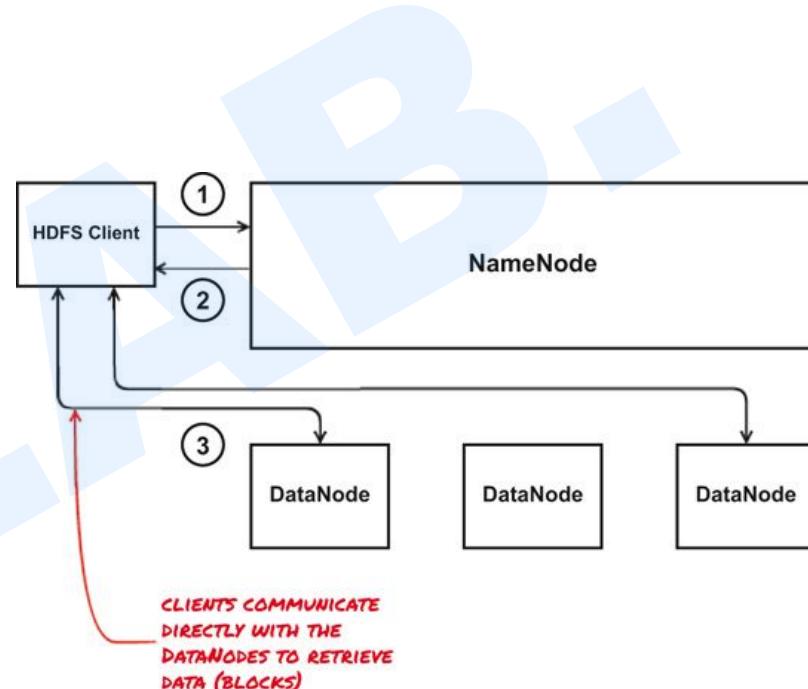
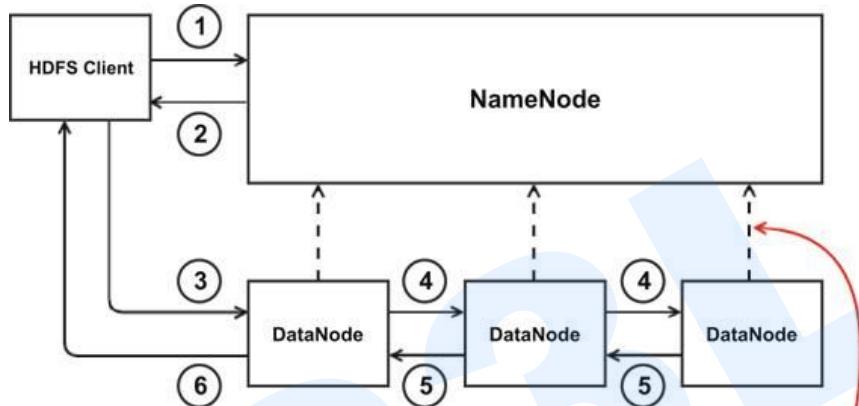
Architecture





HDFS

Write & Read files





HDFS

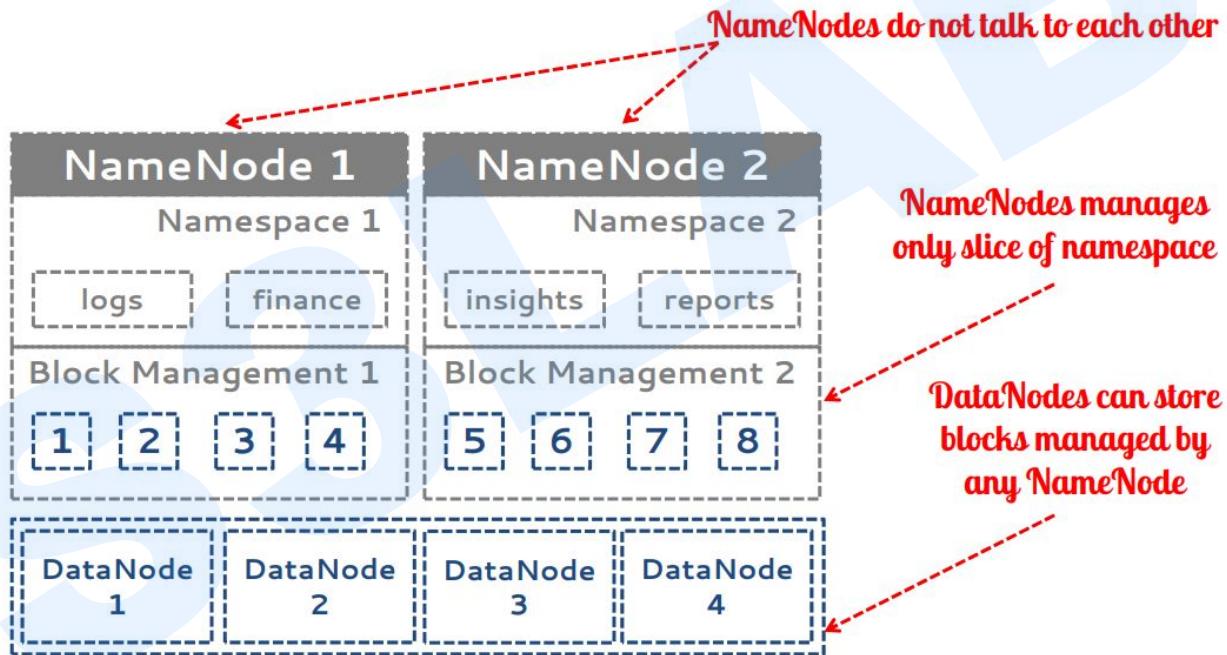
HDFS 2 - Federation

- Removes tight coupling of Block Storage and Namespace
- Scalability & Isolation
- High Availability
- Increased performance



HDFS

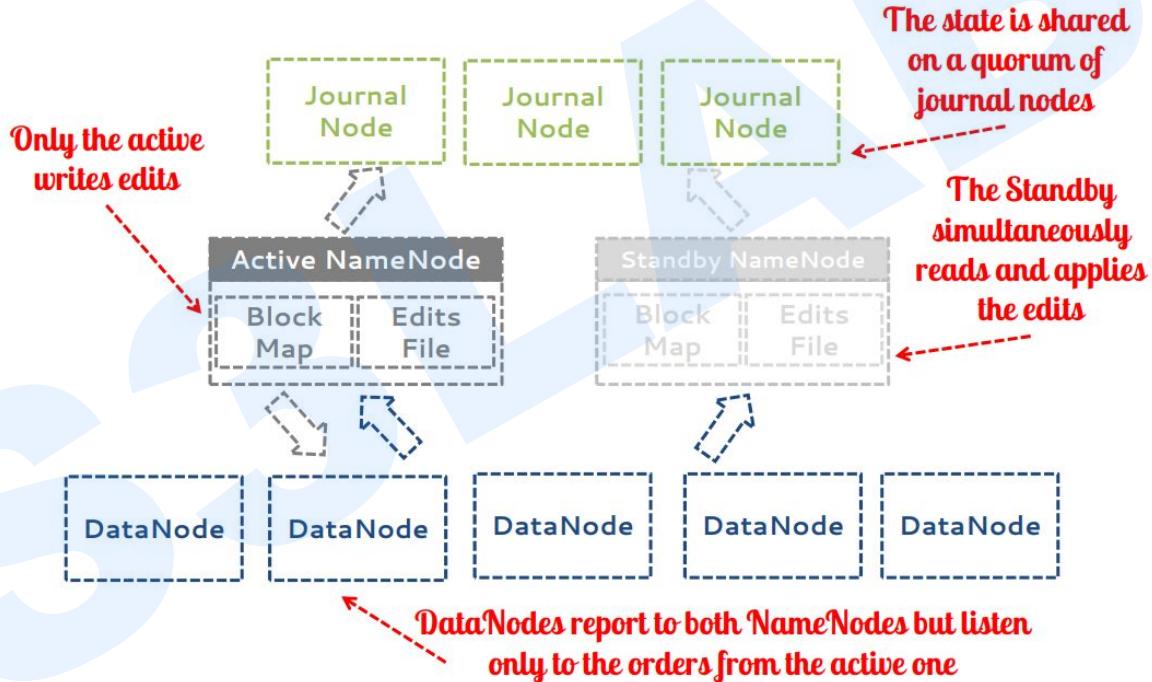
HDFS 2 - Federation





HDFS

HDFS 2 (Federation): Quorum based storage





Map Reduce

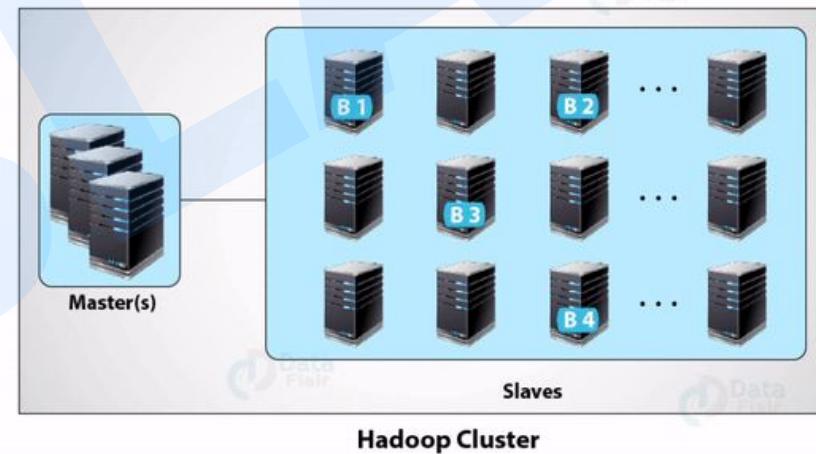
- Programming model for distributed computations at a massive scale
- Execution framework for organizing and performing such computations
- Data locality is king
- JobTracker:
 - takes care of all the job scheduling and assign tasks to Task Trackers.
- TaskTracker
 - a node in the cluster that accepts tasks - Map, Reduce & Shuffle operations - from jobtracker

Map Reduce

flow



How MapReduce works





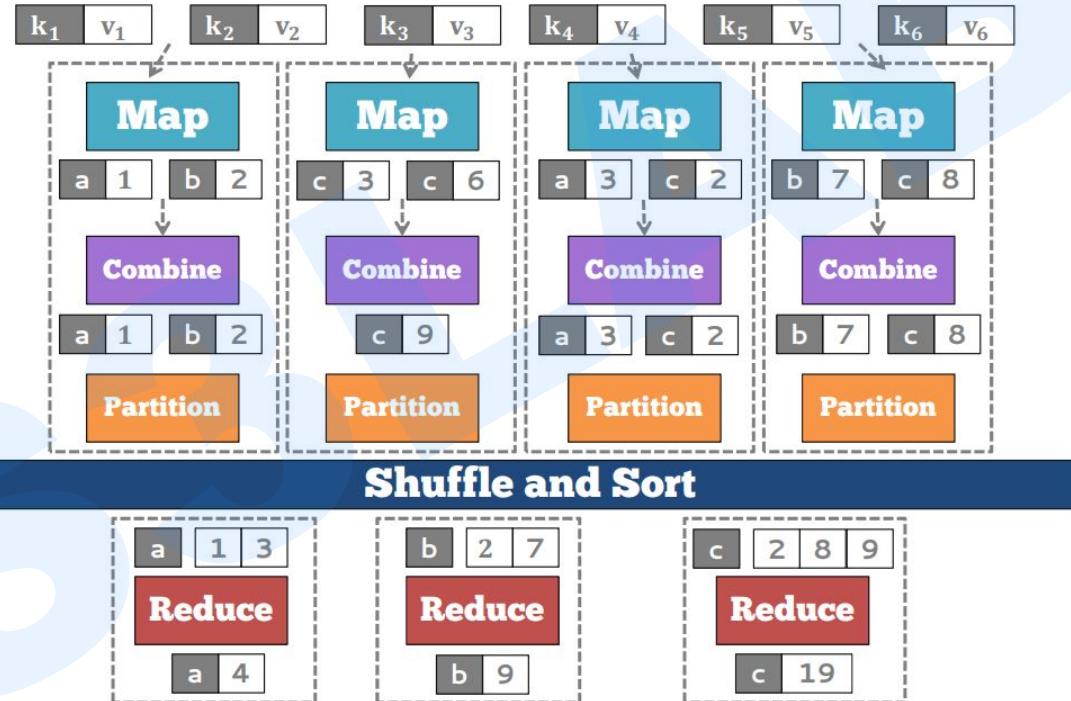
Map Reduce

- The Mapper:
 - Each block is processed in isolation by a map task called mapper
 - Map task runs on the node where the block is stored
 - Iterate over a large number of records
 - Extract something of interest from each
- Shuffle and sort intermediate results
- The Reducer:
 - Consolidate result from different mappers
 - Aggregate intermediate results
 - Produce final output



Map Reduce

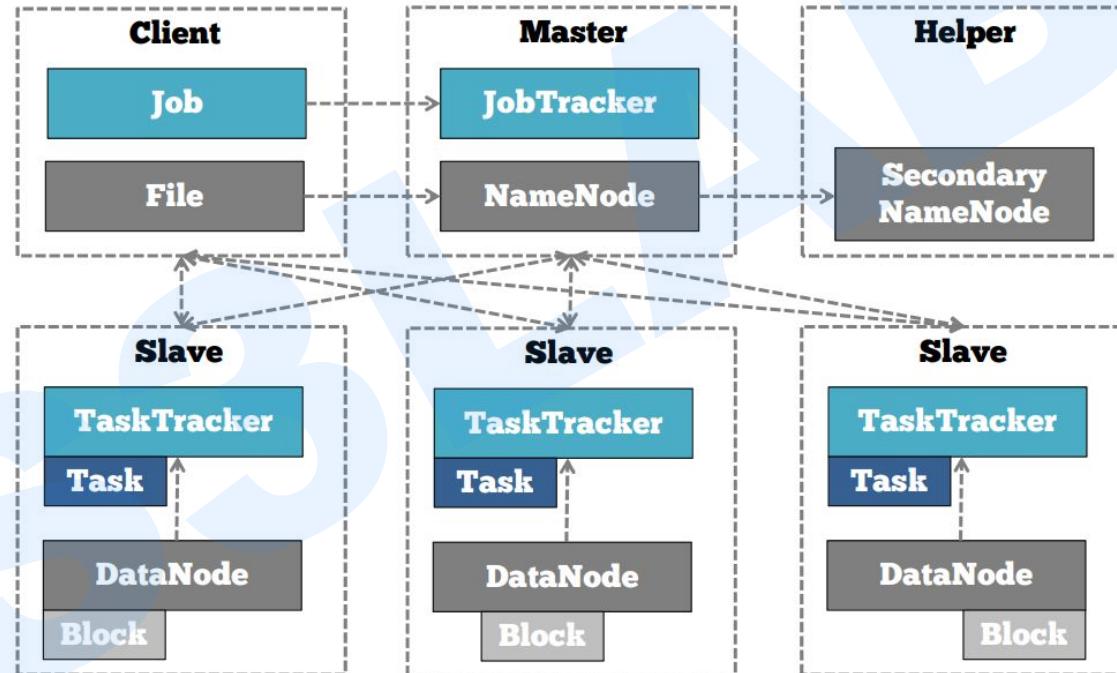
flow





Map Reduce

Combined Hadoop architecture





Map Reduce

Good for ...

- Embarrassingly parallel algorithms
- Summing, grouping, filtering, joining
- Off-line batch jobs on massive data sets
- Analyzing an entire large dataset

SOLVLAB



Map Reduce

OK for ...

- Iterative jobs (i.e., graph algorithms)
 - Each iteration must read/write data to disk
 - I/O and latency cost of an iteration is high

SOLAB



Map Reduce

Not good for ...

- Jobs that need shared state/coordination
 - Tasks are shared-nothing
 - Shared-state requires scalable state store
- Low-latency jobs
- Jobs on small datasets
- Finding individual records



Map Reduce

limitations

- Scalability
 - Maximum cluster size ~ 4,500 nodes, concurrent tasks – 40,000
 - Coarse synchronization in JobTracker
- Availability
 - Failure kills all queued and running jobs
- Hard partition of resources into map & reduce slots
 - Low resource utilization
- Lacks support for alternate paradigms and services
 - Iterative applications implemented using MapReduce are 10x slower



Map Reduce

V2

- Basically a porting to the YARN architecture
- MapReduce becomes a user-land library
- No need to rewrite MapReduce jobs
- Increased scalability & availability
- Better cluster utilization



Hadoop 2.0

history

- Originally conceived & architected by the team at Yahoo!
 - Arun Murthy created the original JIRA in 2008 and now is the YARN release manager
- The team at Hortonworks has been working on YARN for 4 years:
 - 90% of code from Hortonworks & Yahoo!
- Hadoop 2.0 based architecture running at scale at Yahoo!
 - Deployed on 35,000 nodes for 6+ months



Hadoop 2.0

Next-gen platform

Single use system

Batch Apps



Hadoop 1.0

MapReduce

Cluster resource mgmt.
+ data processing

HDFS

Redundant, reliable
storage

Multi-purpose platform

Batch, Interactive, Streaming, ...



Hadoop 2.0

MapReduce

Data processing

Others

Data processing

YARN

Cluster resource management

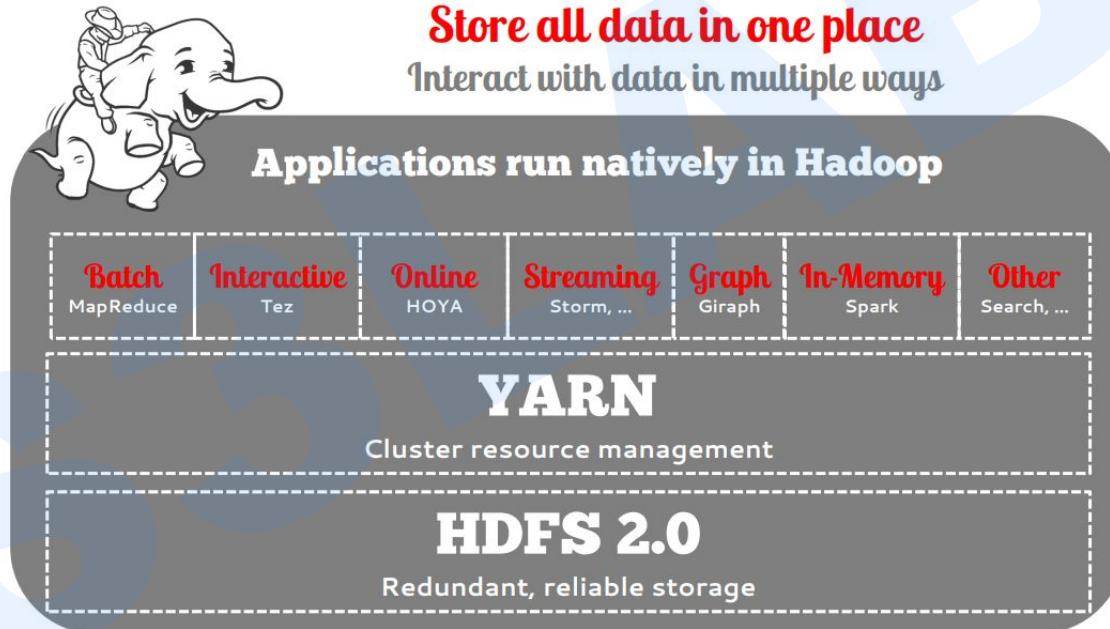
HDFS 2.0

Redundant, reliable storage



Hadoop 2.0

Taking Hadoop beyond batch

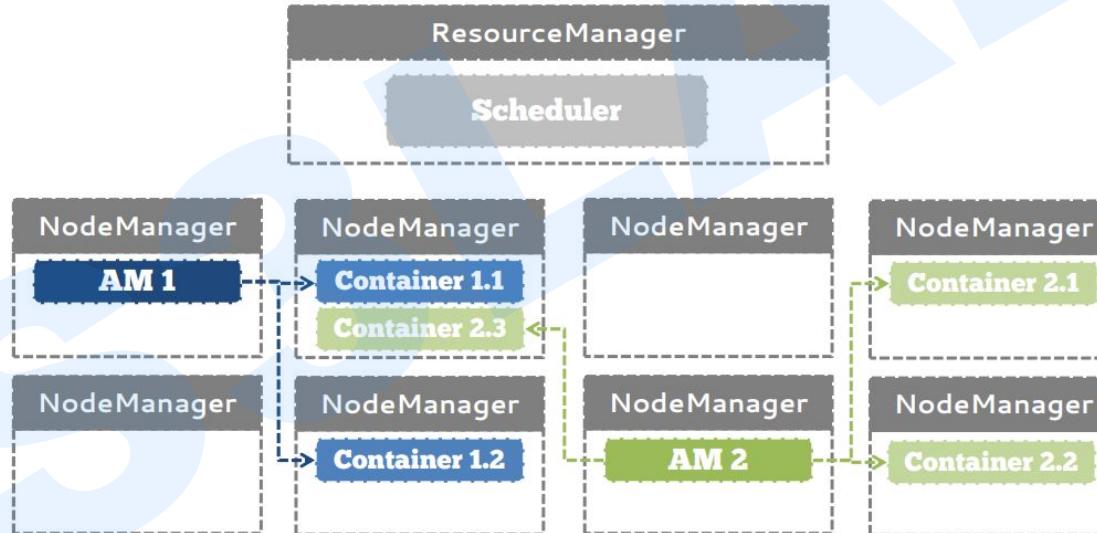




Yarn

Architecture

Split up the two major functions of the JobTracker
Cluster resource management & Application life-cycle management

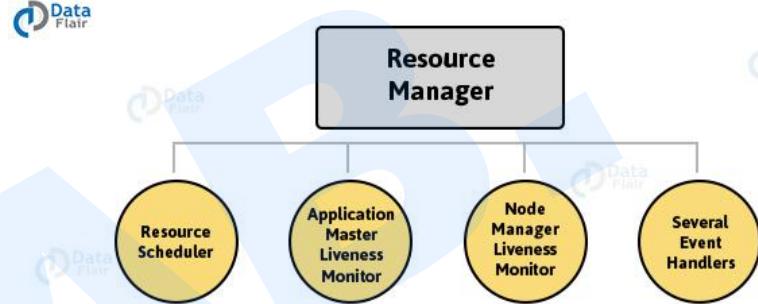


Yarn



Architecture

- **Resource Manager**
 - Resource Manager runs on the master node.
 - It knows where the location of slaves (Rack Awareness).
 - It is aware about how much resources each slave have.
 - Resource Scheduler is one of the important service run by the Resource Manager.
 - Resource Scheduler decides how the resources get assigned to various tasks.
 - Application Manager is one more service run by Resource Manager.
 - Application Manager negotiates the first container for an application.
 - Resource Manager keeps track of the heart beats from the Node Manager.





Yarn

Architecture

- Resource Manager serves an embedded Web UI on port 8088

The screenshot shows the Hadoop YARN Resource Manager Web UI at localhost:8088/cluster/apps. The title bar includes the Hadoop logo and the text "All Applications". The top navigation bar has links for "Cluster Metrics", "About", "Nodes", "Applications", "Scheduler", and "Tools". The "Applications" link is currently selected. The main content area displays a table of running applications:

ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Progress	Tracking UI
application_1412793406652_0002	root	QuasiMonteCarlo	MAPREDUCE	default	Wed, 08 Oct 2014 18:51:01 GMT	N/A	ACCEPTED	UNDEFINED		UNASSIGNED
application_1412793406652_0001	root	QuasiMonteCarlo	MAPREDUCE	default	Wed, 08 Oct 2014 18:45:11 GMT	2014-10-08T18:47:56 GMT	FINISHED	SUCCEEDED	100%	History

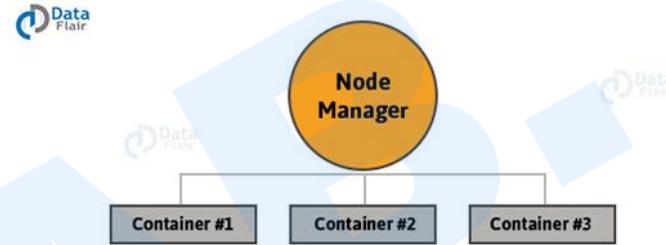
At the bottom, it says "Showing 1 to 2 of 2 entries".



Yarn

Architecture

- Node Manager
 - It runs on slave machines.
 - It manages containers. Containers are nothing but a fraction of Node Manager's resource capacity
 - Node manager monitors resource utilization of each container.
 - It sends heartbeat to Resource Manager.

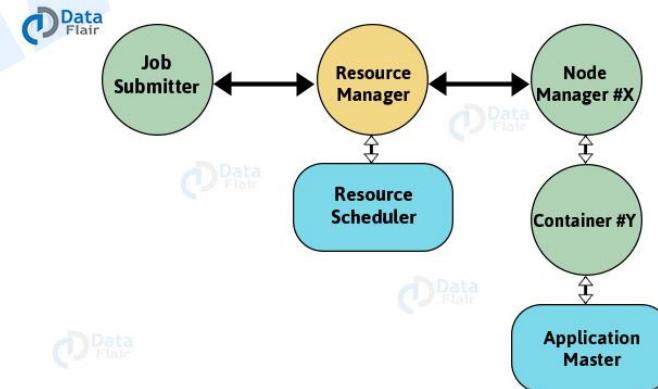




Yarn

Architecture

- Job submitter
 - The client submits the job to Resource Manager.
 - Resource Manager contacts Resource Scheduler and allocates container.
 - Now Resource Manager contacts the relevant Node Manager to launch the container.
 - Container runs Application Master





Yarn

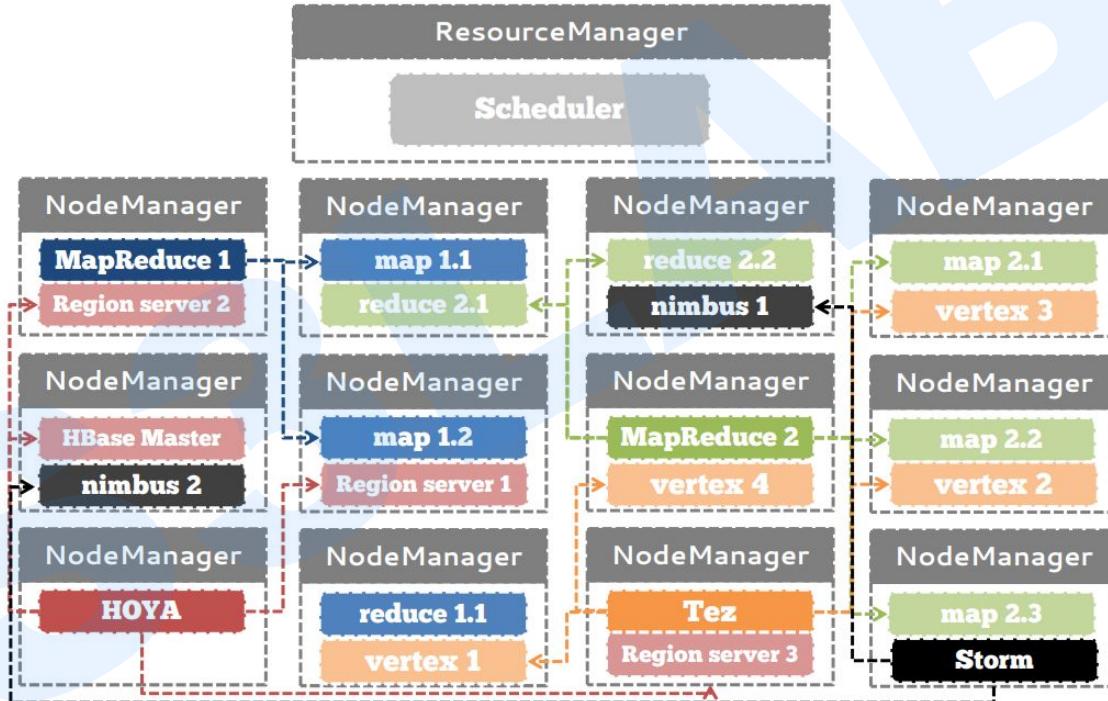
Architecture

- Application Master
 - Per-application
 - Manages application scheduling and task execution
 - e.g. MapReduce Application Master



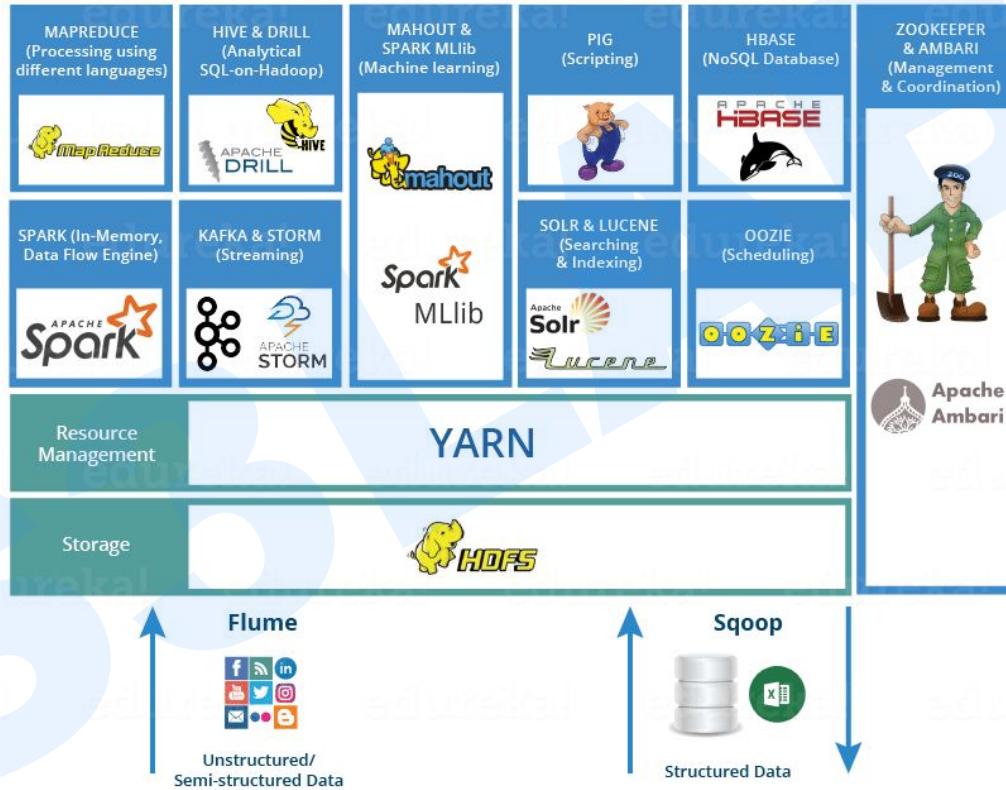
Yarn

Architecture



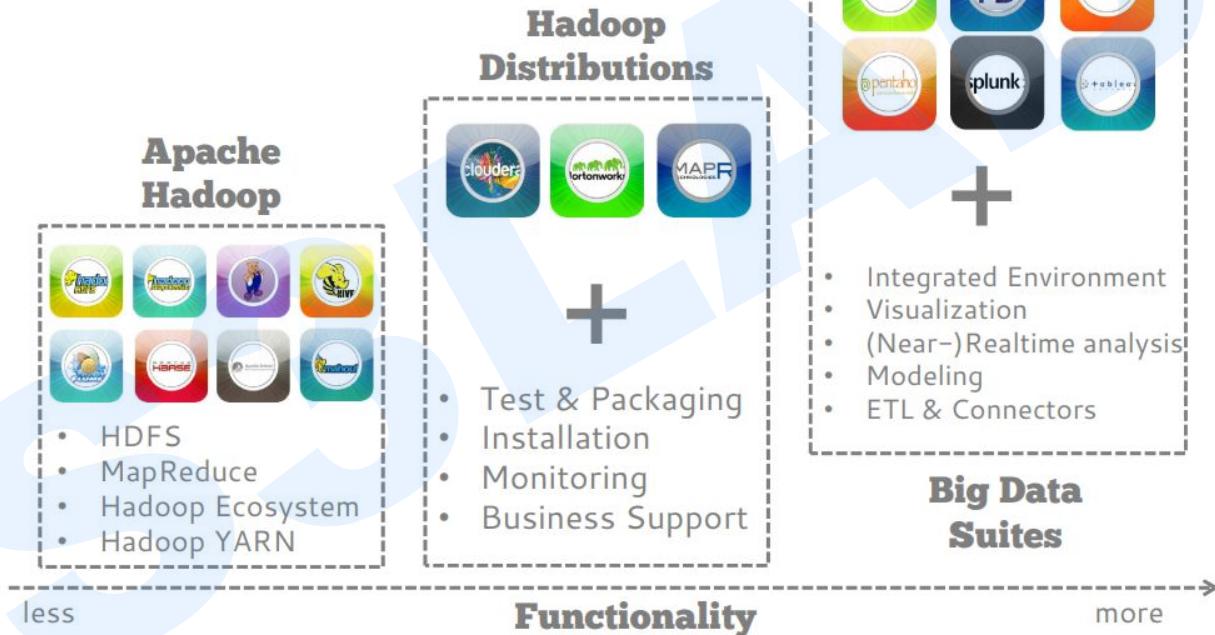


Ecosystem





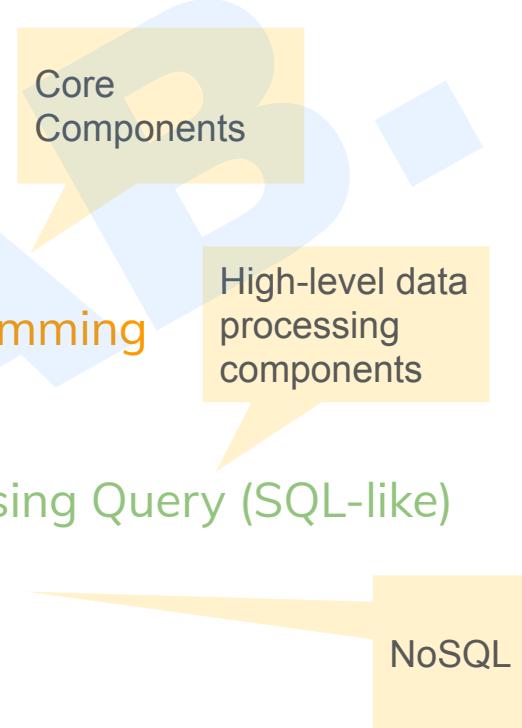
Ecosystem





Ecosystem

- HDFS -> Hadoop Distributed File System
- YARN -> Yet Another Resource Negotiator
- MapReduce -> Data processing using programming
- PIG -> Top-level data processing engine
- HIVE-> Data warehouse on top of hadoop, using Query (SQL-like)
- HBase -> Column oriented NoSQL Database





Ecosystem

- Apache Drill: Schema-free SQL Query Engine for Hadoop
- Solr & Lucene: High-performance text search engine
- Mahout, Spark MLlib -> Data Analysis and machine Learning
- Avro -> data serialization framework
- Thrift -> Interface definition language and binary communication protocol
- Oozie -> Server-based workflow scheduling system
- HCatalog -> Table and storage management layer

Hadoop Data Analysis components

Data serialization

Management Components



Ecosystem

- Flume -> Data collection & aggregation system
- Sqoop -> tool designed for efficiently transferring bulk Data between hadoop and RDBMS
- Chukwa -> Data collection system for monitoring large distributed systems
- Ambari -> Hadoop deployment, management & monitoring tool
- Zookeeper -> Highly reliable distributed coordination system
- Hue (Hadoop user experience) -> Open-source hadoop web interface

Hadoop Data transfer (ingest) components

Monitoring components



Ecosystem

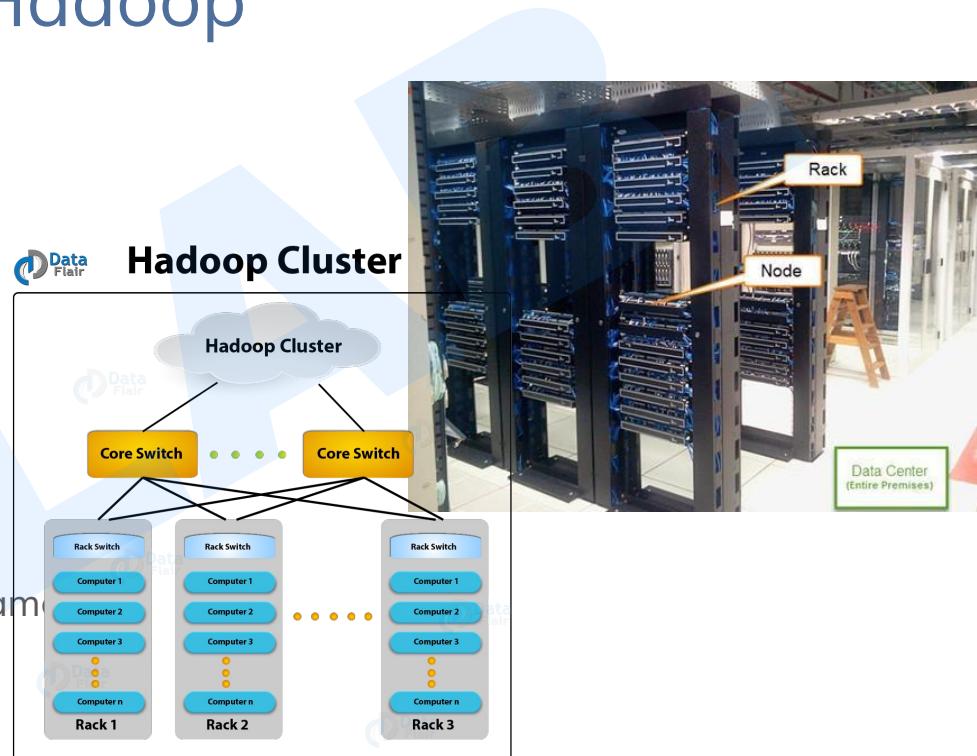
- Spark -> in memory dataflow engine
- Kafka, storm -> support stream data

S3LAB



Network Topology In Hadoop

- Topology of network:
 - Performance
 - Availability & handling of failures
- Hadoop cluster:
 - Processes on the same node
 - Different nodes on the same rack
 - Nodes on different racks of the same data center
 - Nodes in different data centers





When Not to Use Hadoop

Real-time Analytics

- Hadoop works on batch processing.
- Slow processing speed with mapreduce on large data sets.
- Hadoop is not so efficient for iterative processing, as Hadoop does not support cyclic data flow
- Hadoop is not efficient for caching. In Hadoop, MapReduce cannot cache the intermediate data in memory for a further requirement which diminishes the performance of Hadoop.



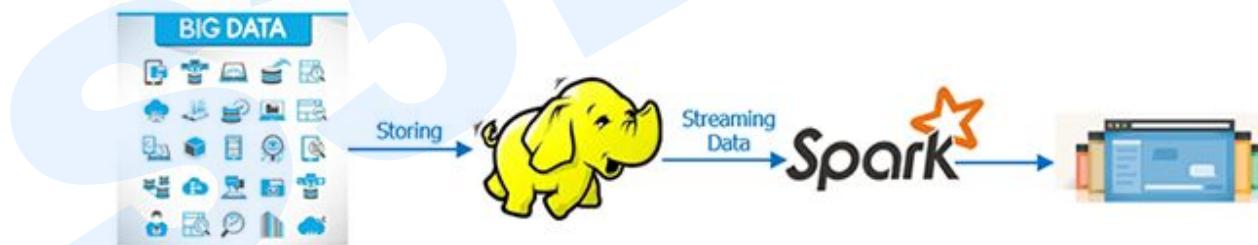
When Not to Use Hadoop

Real-time Analytics

- Solution:
 - Store the big data in HDFS
 - Mount Spark (in-memory processing data) over HDFS (100x faster than mapreduce)
 - Or even using FLINK, it processes faster than spark.

The screenshot shows the Apache Spark Application Overview interface. It displays a yellow elephant icon, the application name 'mywc', and the status 'FINISHED'. Below the status, it says 'Elapsed: 14sec'. To the right, there are two blue clouds with the text '14 sec' and '0.6 sec' respectively. At the bottom, a log output shows the execution details:

```
14/12/08 04:10:05 INFO spark.SparkHadoopWriter: attempt_2014120800410_0005_n_000000_5: Committed
14/12/08 04:10:06 INFO scheduler.TaskSetManager: Finished task 0.0 in stage 5.0 (TID 5) in 296 ms on localhost (1/1)
14/12/08 04:10:06 INFO scheduler.TaskSchedulerImpl: Removed TaskSet 5.0, whose tasks have all completed, from pool
14/12/08 04:10:06 INFO scheduler.DAGScheduler: Stage 5 (saveAsTextFile at <console>:14) finished in 0.279 s
14/12/08 04:10:06 INFO spark.SparkContext: Job finished: saveAsTextFile at <console>:14, took 0.626043309 s
14/12/08 04:10:06 INFO executor.Executor: Finished task 0.0 in stage 5.0 (TID 5). 826 bytes result sent to driver
wordCounts: Unit = ()
```

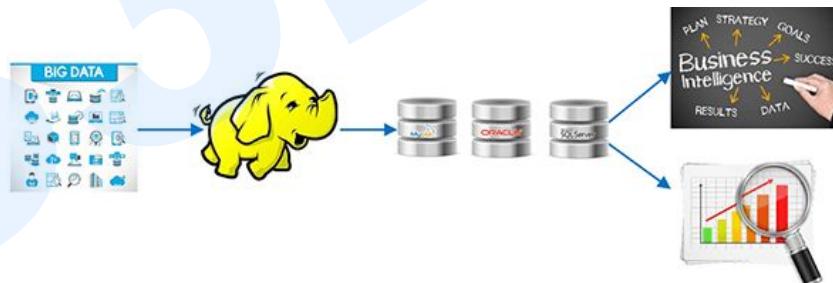




When Not to Use Hadoop

Not a Replacement for Existing Infrastructure

- Hadoop is not a replacement for your existing data processing infrastructure. However, you can use Hadoop along with it.
- Solution:
 - Hadoop is not going to replace your database, but your database isn't likely to replace Hadoop either.
 - Different tools for different jobs, as simple as that.

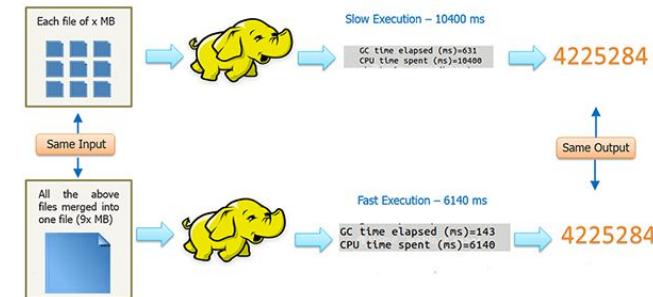




When Not to Use Hadoop

Multiple Smaller Datasets

- Hadoop is not recommended for small-structured datasets, small files (< default block size: 128MB). For a small data analytics, Hadoop can be costlier than other tools.
- Solution:
 - Merge the small file to create bigger files and then copy bigger files to HDFS





When Not to Use Hadoop

Multiple Smaller Datasets

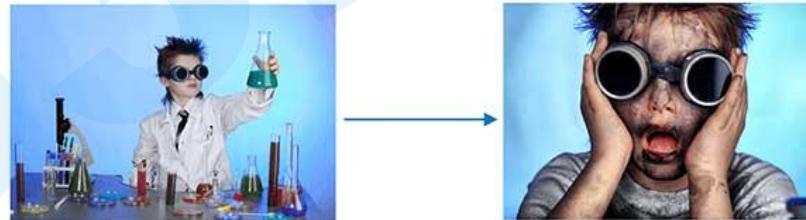
- Solution:
 - Using HAR files (Hadoop archives)
 - Sequence files: we use filename -> key, file contents -> value. Put them into a single sequence file with a program for files and process them in a streaming fashion operating on the sequence file. MapReduce can break the sequence file into chunks and operate on each chunk independently because the sequence file is splittable.
 - Storing files in HBase. We are not actually storing millions of small files into HBase, rather adding the binary content of the file to a cell.



When Not to Use Hadoop

Novice Hadoopers

- Hadoop is a technology which should come with a disclaimer: “Handle with care”. You should know it before you use it or else you will end up like the kid below.
- Not easy to use, does not have any type of abstraction, mapreduce has no interactive mode.





When Not to Use Hadoop

Novice Hadoopers

- Solutions:

- Using Spark with RDD (Resilient Distributed DataSet) abstraction for the batch or Flink has dataset abstraction.
- Spark has interactive mode, Flink has high-level operators.



When Not to Use Hadoop

Where Security is the primary Concern?

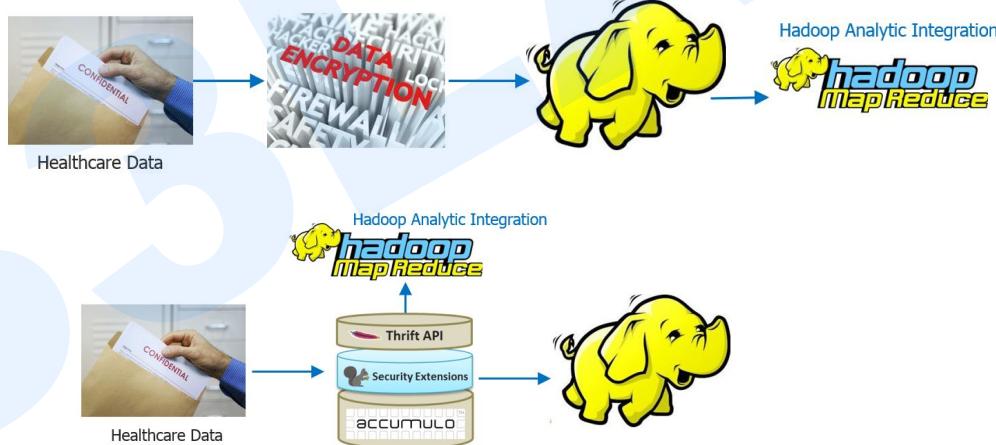
- Many enterprises — especially within highly regulated industries dealing with sensitive data — aren't able to move as quickly as they would like towards implementing Big Data projects and Hadoop.
- Hadoop is missing encryption.
- Supports Kerberos authentication, which is hard to manage.



When Not to Use Hadoop

Security

- Solution:
 - Using Spark, even spark also can use HDFS ACLs and file-level permissions in HDFS.
 - Spark can run on yarn -> Kerberos authentication





NoSQL Database

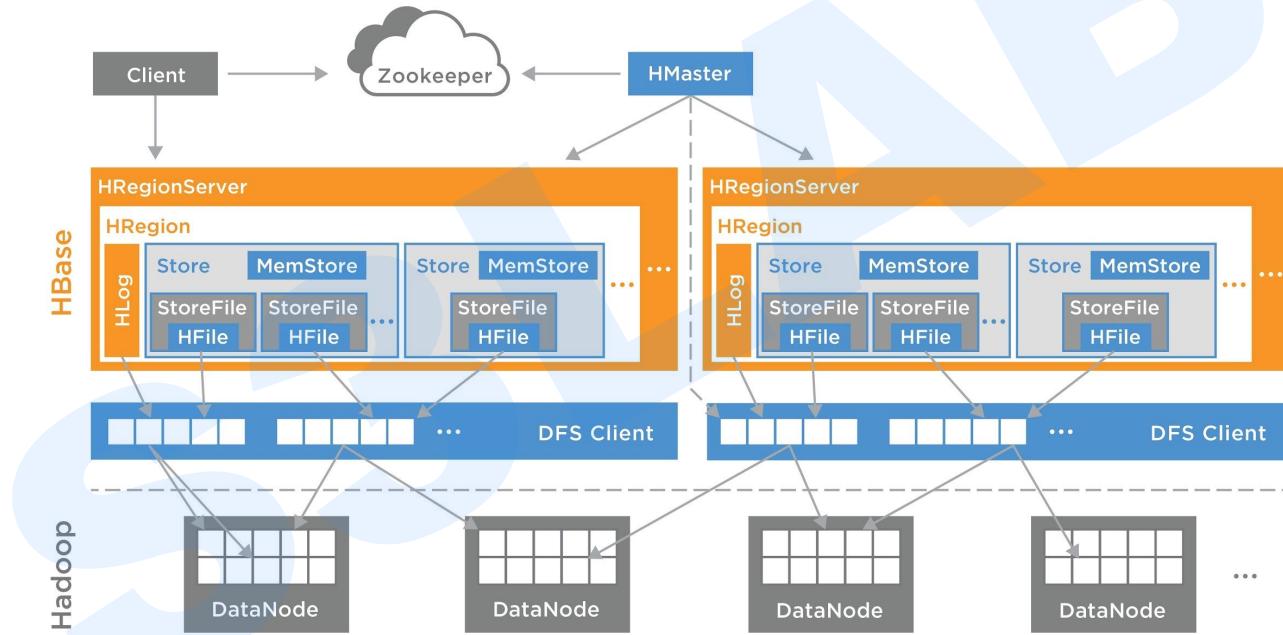
HBase

- HBase is an open source, non-relational, distributed database modeled after Google's BigTable.
- It runs on top of Hadoop and HDFS, providing BigTable-like capabilities for Hadoop.



NoSQL Database

HBase





NoSQL Database

HBase - Features

- Type of NoSql database (column oriented)
- Strongly consistent read and write
- Automatic sharding
- Automatic RegionServer failover
- Hadoop / HDFS Integration



NoSQL Database

HBase - Features

- HBase supports massively parallelized processing via MapReduce for using HBase as both source and sink.
- HBase supports an easy to use Java API for programmatic access.
- HBase also supports Thrift and REST for non-Java front-ends.



NoSQL Database

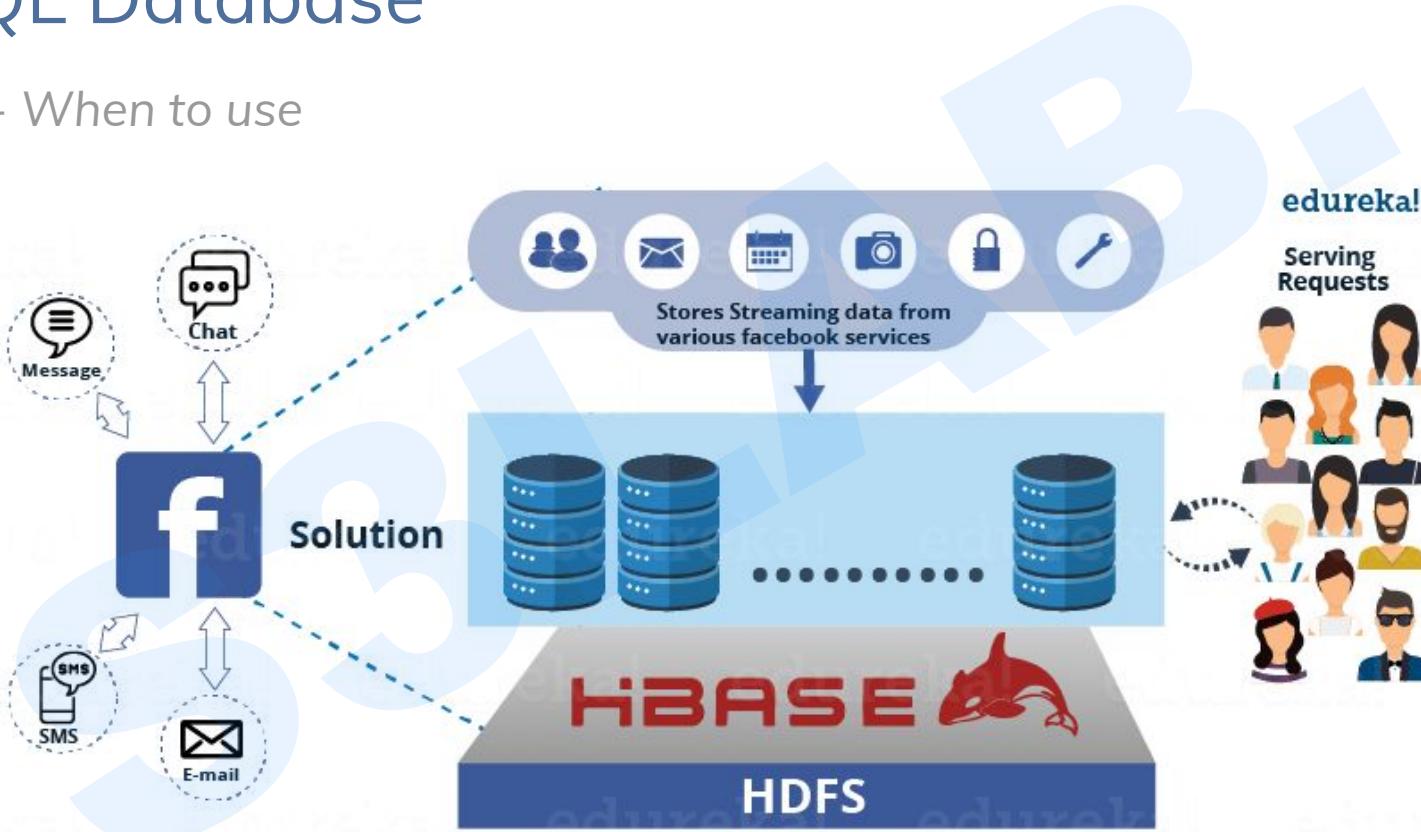
HBase - When to use

- When there is real big data: millions or billions of rows, in other way data can not store in a single node.
- When random read/write access to big data
- When require to do thousands of operations on big data
- When there is no need of extra features of RDMS like typed columns, secondary indexes, transactions, advanced query languages, etc.
- When there is enough hardware.



NoSQL Database

HBase - When to use





NoSQL Database

Difference between **Hbase** and HDFS

HDFS	Hbase
Good for storing large file	Built on top of HDFS. Good for hosting very large tables like billions of rows X millions of column
Write once. Append to files in some of recent versions but not commonly used	Read/write many
No random read/write	Random read/write
No individual record lookup rather read all data	Fast records lookup(update)



NoSQL Database

Hoya: HBase on Yarn

- Create on-demand HBase clusters
- Configure different HBase instances differently
- Better isolation
- Create (transient) HBase clusters from MapReduce jobs
- Elasticity of clusters for analytic / batch workload processing
- Better cluster resources utilization



High-Level Data Process Components

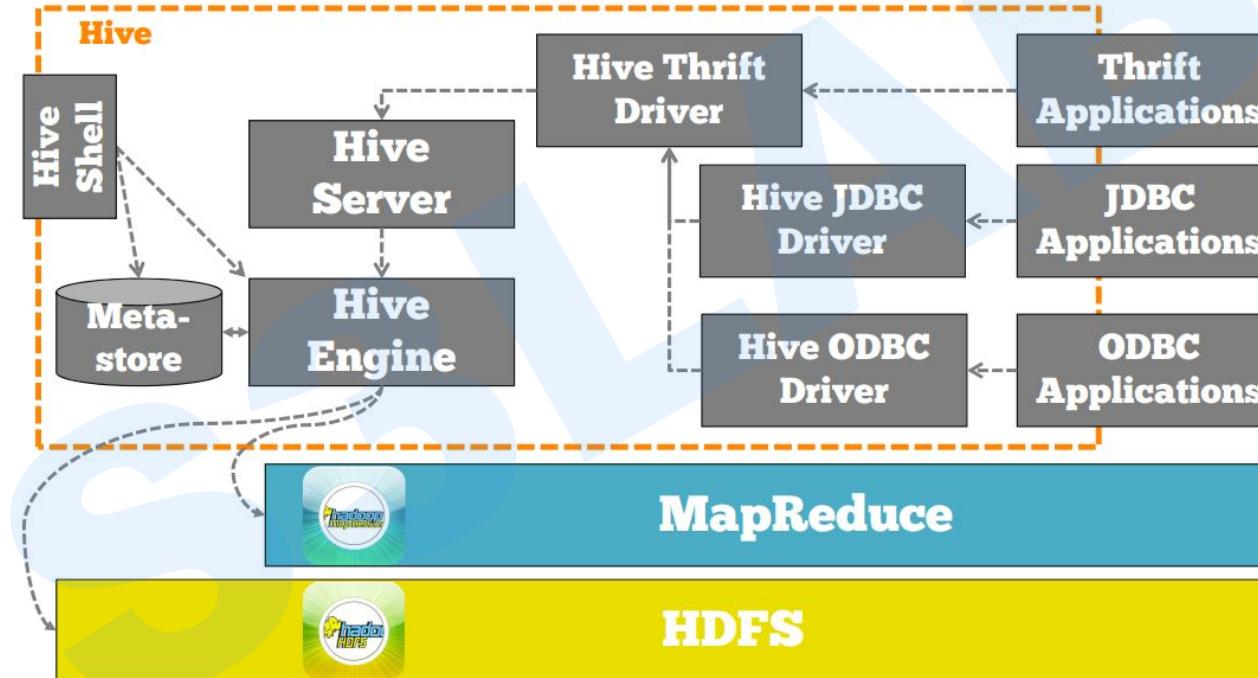
Hive

- An sql like interface to Hadoop.
- Data warehouse infrastructure built on top of Hadoop
- Provide data summarization, query and analysis
- Query execution via MapReduce
- Hive interpreter convert the query to Map-reduce format.
- Open source project.
- Developed by Facebook
- Also used by Netflix, Cnet, Digg, eHarmony etc.



High-Level Data Process Components

Hive - architecture





High-Level Data Process Components

Hive

- HiveQL example:

```
SELECT customerId, max(total_cost) from hive_purchases GROUP BY  
customerId HAVING count(*) > 3;
```



High-Level Data Process Components

Pig

- A scripting platform for processing and analyzing large data sets
- Apache Pig allows to write complex MapReduce programs using a simple scripting language.
- Made of two components:
 - High level language: Pig Latin (data flow language).
 - Pig translates Pig Latin script into MapReduce to execute within Hadoop.
- Open source project
- Developed by Yahoo



High-Level Data Process Components

Pig

- Pig Latin example:

```
A = LOAD 'student' USING PigStorage() AS (name:chararray, age:int,  
gpa:float);
```

```
X = FOREACH A GENERATE name,$2;
```

```
DUMP X;
```



High-Level Data Process Components

Hive & Pig

- Both require compiler to generate Mapreduce jobs
- Hence high latency queries when used for real time responses to ad-hoc queries
- Both are good for batch processing and ETL jobs
- Fault tolerant



High-Level Data Process Components

Impala

- Cloudera Impala is a query engine that runs on Apache Hadoop.
- Similar to HiveQL.
- Does not use Map-reduce
- Optimized for low latency queries
- Open source apache project
- Developed by Cloudera
- Much faster than Hive or pig



High-Level Data Process Components

Impala, Pig and Hive

Description of Feature	Pig	Hive	Impala
SQL based query language	No	yes	yes
Schema	optional	required	required
Process data with external scripts	yes	yes	no
Extensible file format support	yes	yes	no
Query speed	slow	slow	fast
Accessible via ODBC/JDBC	no	yes	yes



Data transfer components

Sqoop (Sql-to-hadoop)

- Command-line interface for transforming data between RDBMS & Hadoop
- Parallelized data transfer with MapReduce
- Support incremental imports
- Imports use to populate tables in Hadoop
- Exports use to put data from Hadoop into relational database
- Sqoop2 -> Sqoop-as-a-Service: server-based implementation of Sqoop





Data transfer components

Sqoop - How it works

- The dataset being transferred is broken into small blocks.
- Map only job is launched.
- Individual mapper is responsible for transferring a block of the dataset.



Data transfer components

Sqoop - How it works

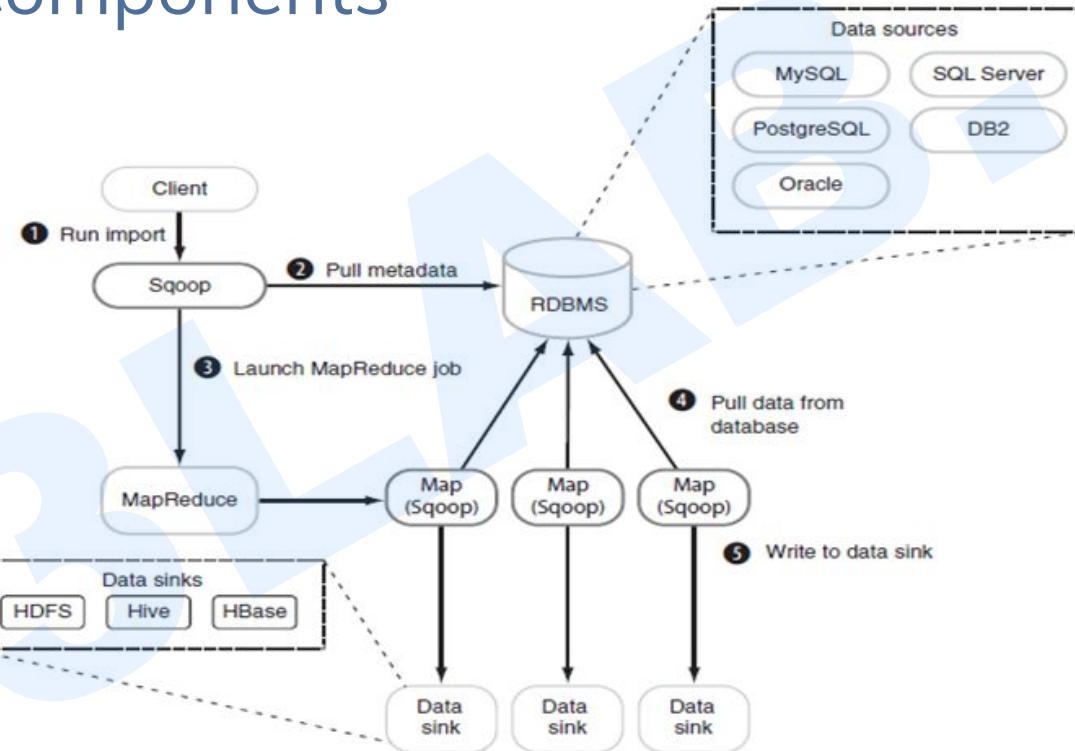


Figure 2.20 Five-stage Sqoop import overview: connecting to the data source and using MapReduce to write to a data sink



Data transfer components

Flume

- Apache Flume is a distributed, reliable, and available service for efficiently collecting, aggregating, and moving large amounts of streaming data (ingest) into the Hadoop Distributed File System (HDFS).



Data transfer components

Flume - How it works

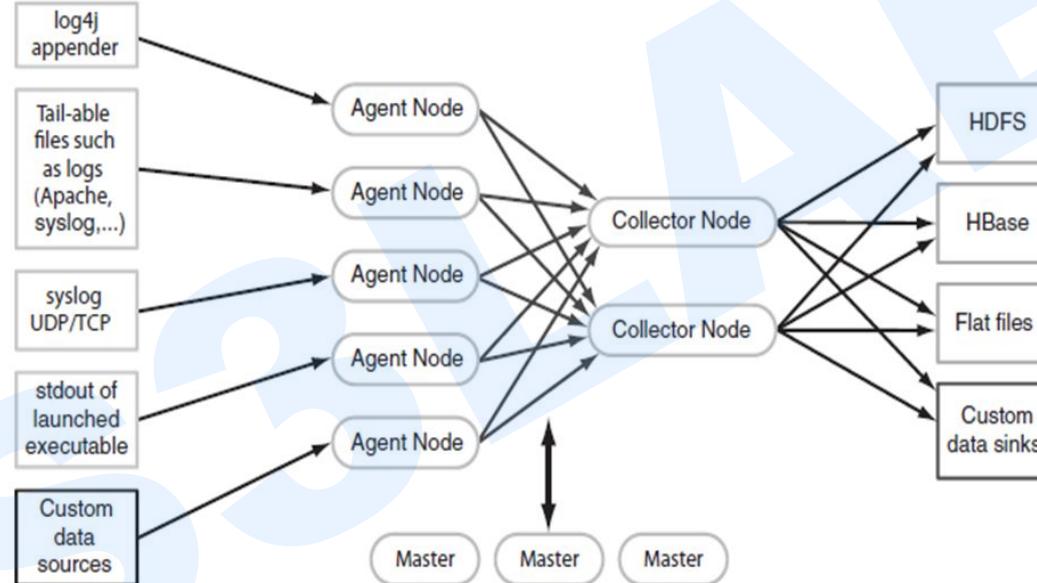


Figure 2.2 Flume architecture for collecting streaming data



Data transfer components

Flume - How it works

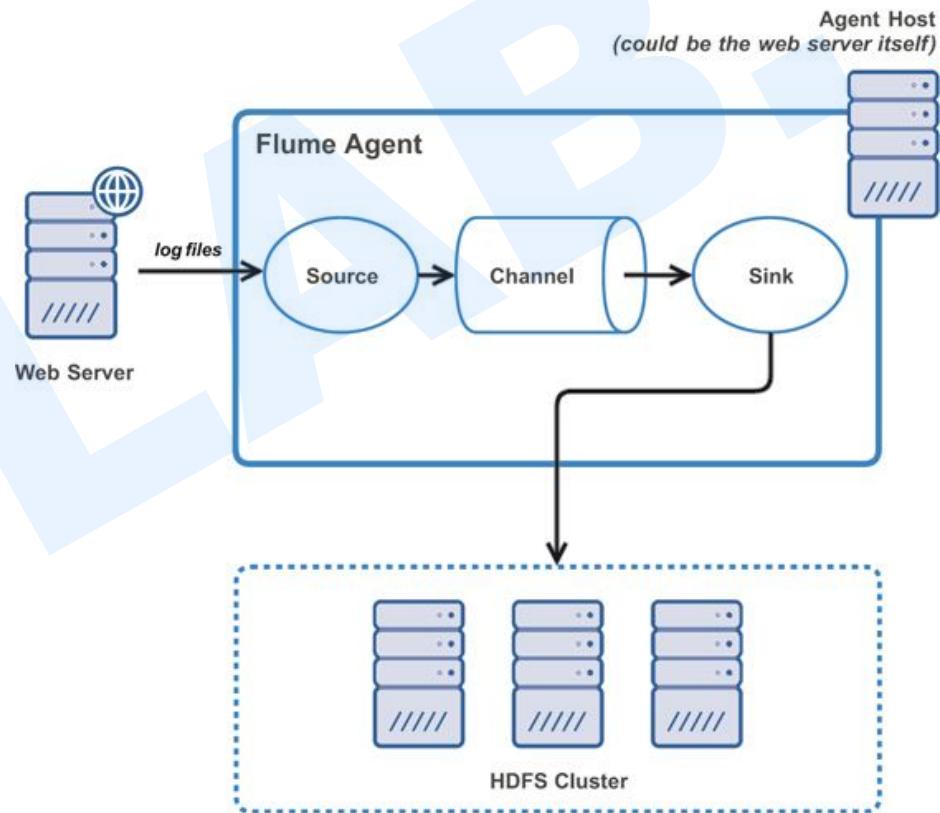
- Data flows like
- Agent tier -> Collector tier -> Storage tier
- Agent nodes are typically installed on the machines that generate the logs and are data's initial point of contact with Flume. They forward data to the next tier of collector nodes, which aggregate the separate data flows and forward them to the final storage tier.



Data transfer components

Flume - Agent architecture

- Sources:
 - HTTP, Syslog, JMS, Kafka, Avro, Twitter - stream api for tweets download, ...
- Sink:
 - HDFS, Hive, HBase, Kafka, Solr, ...
- Channel:
 - File, JDBC, Kafka, ...

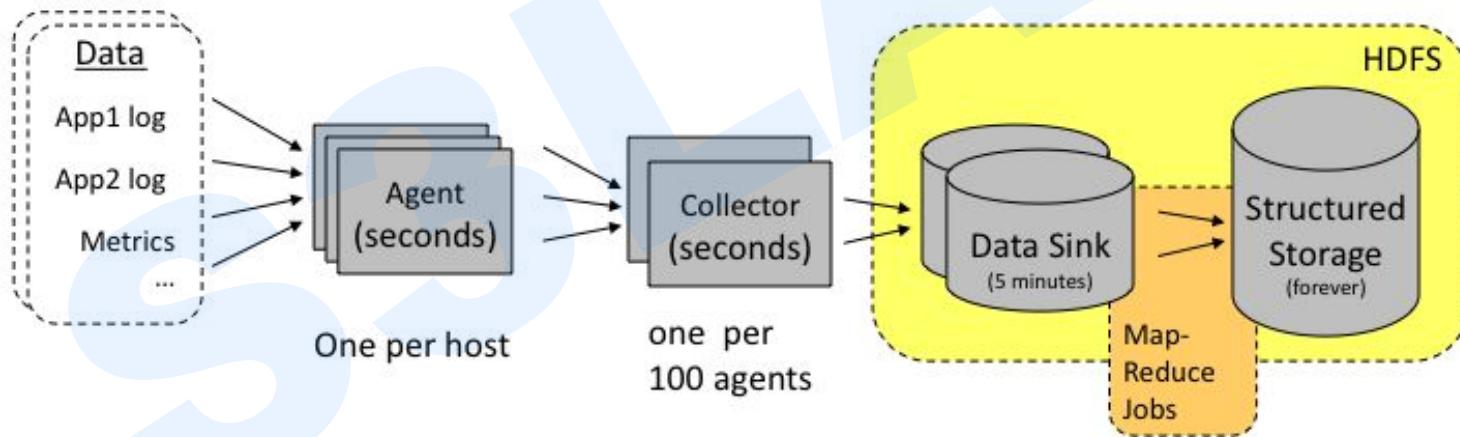




Data transfer components

Chukwa

- a scalable system for collecting logs and other monitoring data and processing the data with MapReduce





Data transfer components

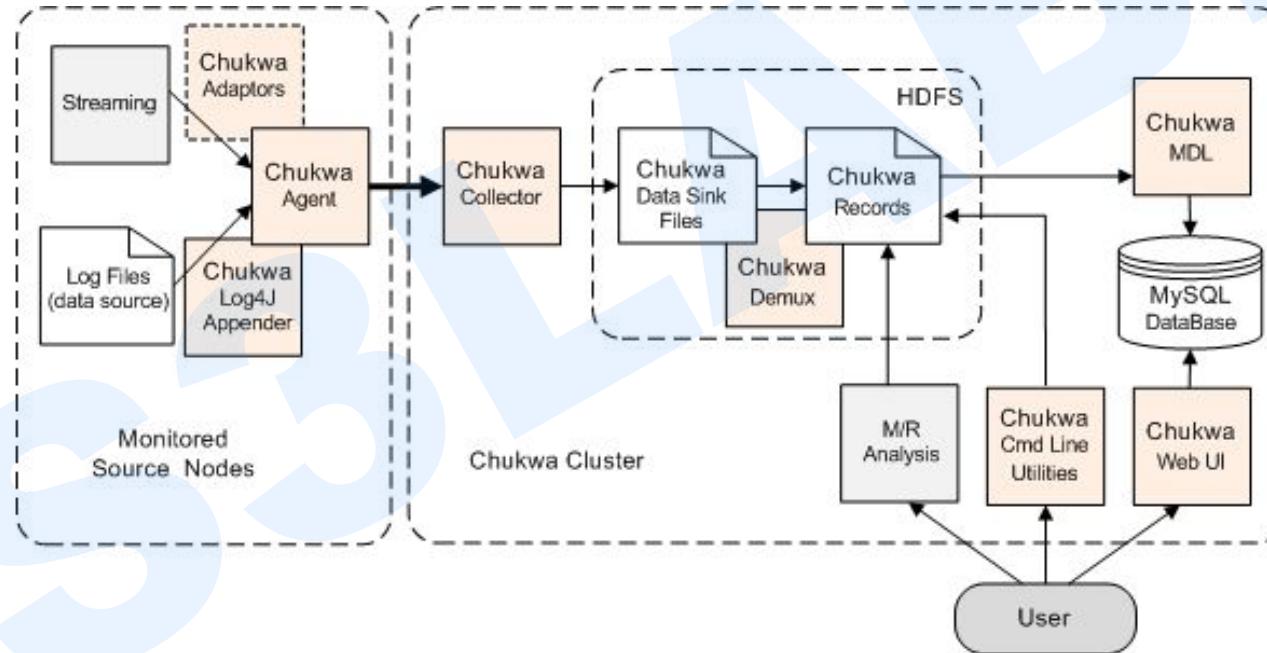
Chukwa

- **Agents** that run on each machine and emit data.
 - Adaptor: dynamically controllable data sources. Emit data in **Chunks** (a sequence of bytes, with some metadata).
- **Collectors** that receive data from the agent and write it to stable storage.
- **MapReduce jobs** for parsing and archiving the data.
- **HICC**, the Hadoop Infrastructure Care Center; a web-portal style interface for displaying data.



Data transfer components

Chukwa - HICC





Data Serialization Components

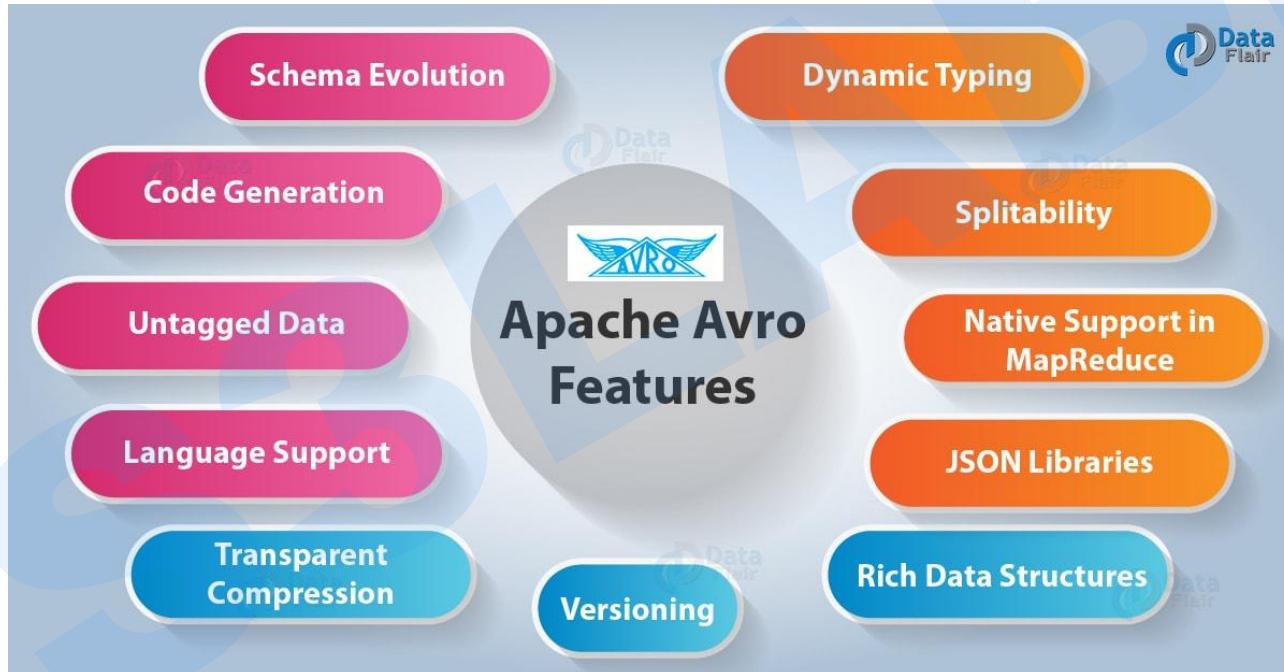
AVRO

- A language-neutral data serialization system.
- Avro uses JSON based schemas
- Uses RPC calls to send data
- During the data exchange, schema's sent.



Data Serialization Components

Avro - features





Data Serialization Components

Avro - General working

- **Step 1 – Create schemas.** Here you need to design Avro schema according to your data.
- **Step 2 – Read the schemas into your program. It is done in two ways –**
 - **By Generating a Class Corresponding to Schema** – Compile the schema using Avro. This generates a class file corresponding to the schema
 - **By Using Parsers Library** – You can directly read the schema using parsers library.



Data Serialization Components

Avro - General working

- **Step 3** – Serialize the data using the serialization API provided for Avro, which is found in the **package org.apache.avro.specific**.
- **Step 4** – Deserialize the data using deserialization API provided for Avro, which is found in the **package org.apache.avro.specific**.



Data Serialization Components

Thrift

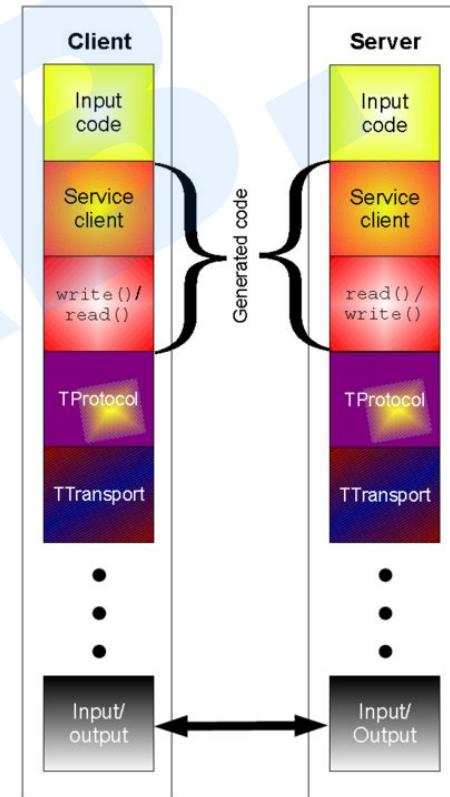
- Define data types and service interface with IDL (interface definition language)
- The IDL content are stored as `.thrift`.
- Thrift provides clean abstractions for data transport, data serialization, and application level processing.



Data Serialization Components

Thrift

- Apache Thrift is a set of code-generation tools that allows developers to build RPC clients and servers by just defining the data types and service interfaces in a simple definition file.
Given this file as an input, code is generated to build RPC clients and servers that communicate seamlessly across programming languages





Data Serialization Components

Thrift

Technologies	Protobuf	Thrift	Avro
Created	2001 (2008)	2007	2009
Creator / Maintainer	Google / Google	Facebook / Apache	Doug cutting / Apache
Schema evolution	Field Tag	Field Tag	Schema
Static/Dynamic	Yes/No	Yes/No	Yes/Yes
Hadoop support	No	No	Yes
RPC	No	Yes	Yes
Used by	Google	Facebook, Cassandra	Hadoop, Liveperson
Lang support	Good	Great	Good



Monitoring Components

Hue

- Graphical front end to the cluster.
- Open source web interface.
- Makes Hadoop platform (HDFS, yarn, Solr, Pig, Map reduce, oozie, Hive, sqoop, impala, etc.) easy to use



Monitoring Components

Hue

Screenshot of the Hue interface showing a query editor and a chart.

Query Editor:

```
Sample: Salary growth  Salary growth (sorted) from 2007-08

1 SELECT s07.description, s07.salary, s08.salary,
2     s08.salary - s07.salary
3 FROM
4     sample_07 s07 JOIN sample_08 s08
5     ON s07.code = s08.code
6 WHERE
7     s07.salary < s08.salary
8 ORDER BY s08.salary-s07.salary DESC
9 LIMIT 20
```

Chart:

Bar chart titled "Salary growth (sorted) from 2007-08" showing salary differences across various professions.

Profession	Salary Difference
Dentists, all Surgeons other specialists	~120,000
Oral and maxillofacial surgeons	~180,000
Natural managers	~100,000
Physicians and surgeons, all other	~150,000
Orthodontists	~180,000
Internists, general	~160,000
Political scientists	~80,000
Obstetricians and gynecologists	~180,000
Chief executives	~150,000
Rotary drill operators, oil and gas	~40,000
Pediatricians	~140,000
Biochemists and biophysicists	~70,000
Family and general practitioners	~150,000
Medical scientists, except competitors	~70,000
Athletes	~70,000
Animal scientists	~50,000
Dentists, general	~150,000
Education psychologists, administrative and postsecondary	~80,000
Psychologists, all other	~80,000



Monitoring Components

ZooKeeper

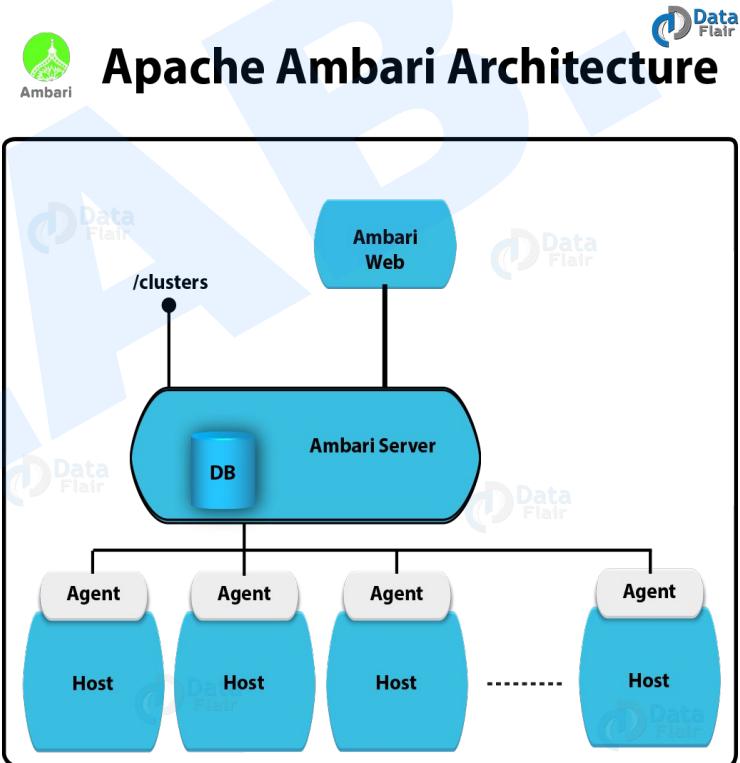
- Because coordinating distributed systems is a Zoo.
- ZooKeeper is a centralized service for maintaining configuration information, naming, providing distributed synchronization, and providing group services.



Monitoring Components

Ambari

- an open source web-based management tool which manages, monitors as well as provisions the health of Hadoop clusters

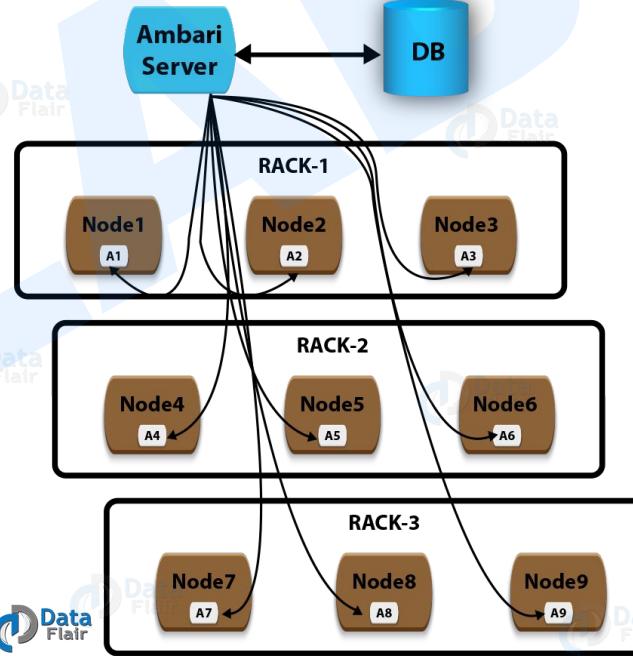




Monitoring Components

Ambari

Working of Apache Ambari





Data Analysis Components

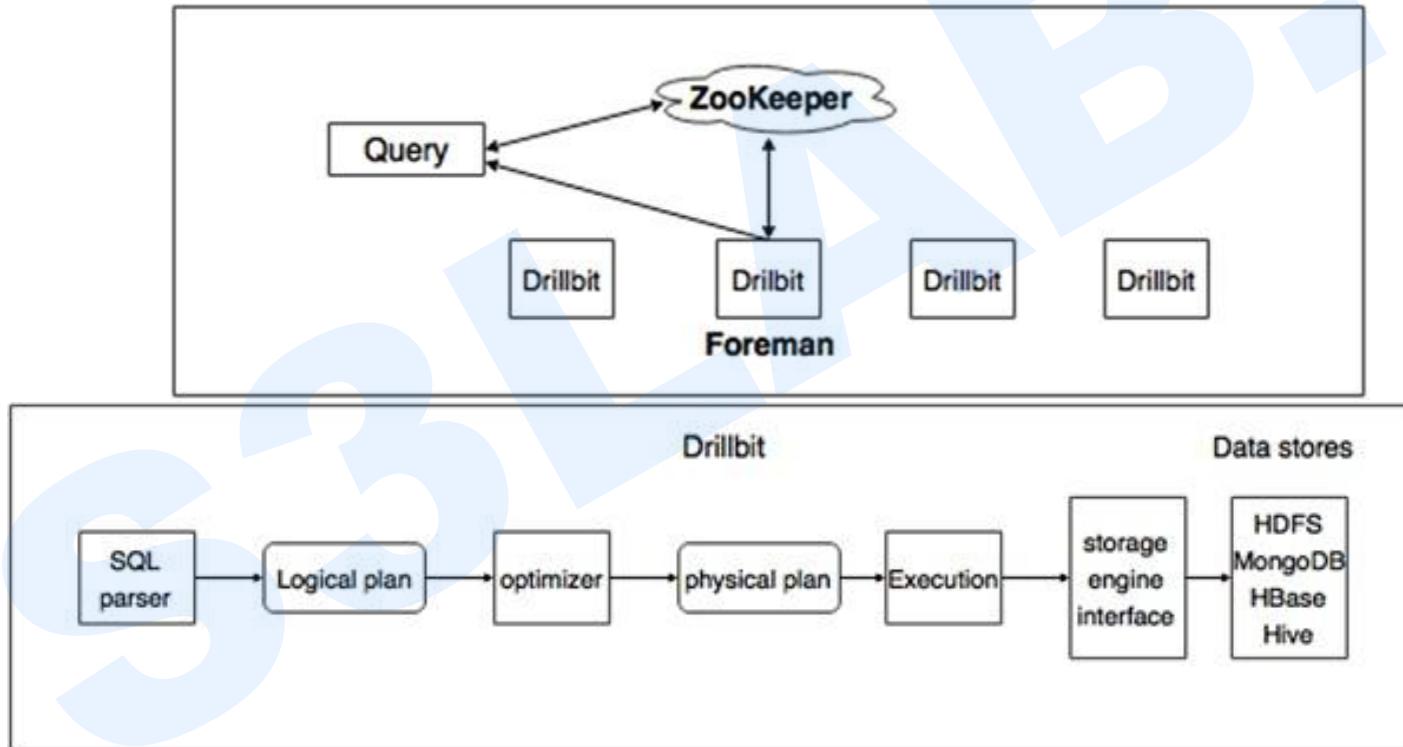
Drill

- A low latency schema-free query engine for big data
- Users can query the data using a standard SQL and BI Tools, which doesn't require to create and manage schemas
- Drill uses a JSON document model internally which allows it to query data of any structure.
- Drill works with a variety of non-relational data stores, including Hadoop, NoSQL databases (MongoDB, HBase), Local files, NAS and cloud storage like Amazon S3, Azure Blob Storage, etc



Data Analysis Components

Drill





Data Analysis Components

Mahout

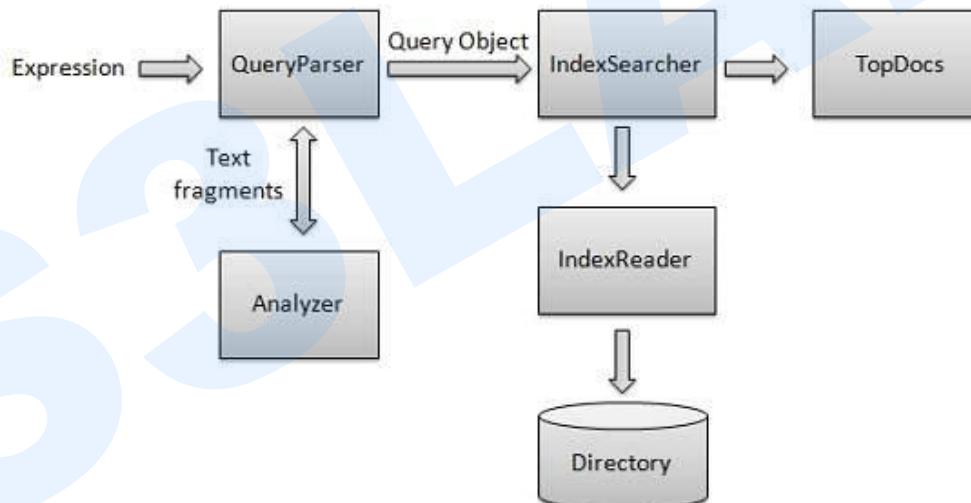
- A mahout is one who drives an elephant as its master.
- is an open source project that is primarily used for creating scalable machine learning algorithms, implemented on top of Apache Hadoop® and using the MapReduce paradigm. It implements popular machine learning techniques such as:
 - Recommendation
 - Classification
 - Clustering



Data Analysis Components

Lucene

- Lucene is an open-source Java full-text search library which makes it easy to add search functionality to an application or website.





Management Components

Oozies

- Workflow scheduler to manage hadoop and related jobs
- Developed first in Bangalore by Yahoo
- DAG(Direct Acyclic Graph): Acyclic means a graph cannot have any loops and action members of the graph provide control dependency.
Control dependency means a second job cannot run until a first action is completed
- Oozie definitions are written in hadoop process definition language (hPDL) and coded as an xml file (WORKFLOW.XML)



Management Components

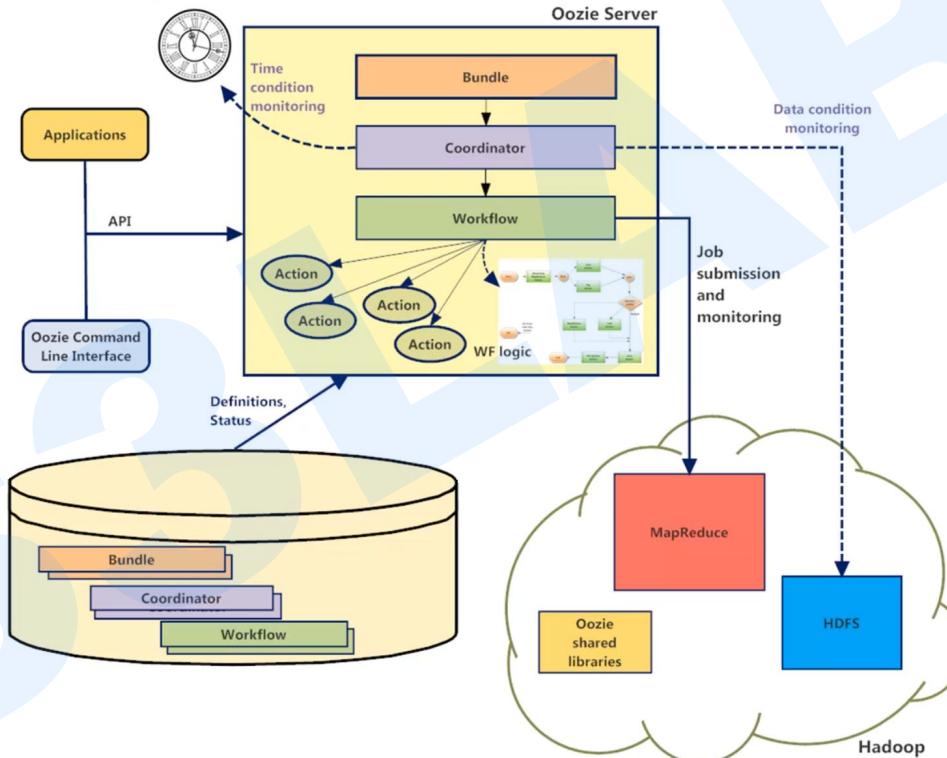
Oozies

- Workflow contains:
 - Control flow nodes (defines start, end and execution path of the workflow):
START, FORK, JOIN, DECISION, KILL, END
 - Action nodes (trigger execution of tasks): Java MapReduce, Streaming
MapReduce, Pig, Hive, Sqoop, FileSystem tasks, Distributed copy, Java programs,
Shell scripts, Http, Email, Oozie sub workflows, ...



Management Components

Oozies





Management Components

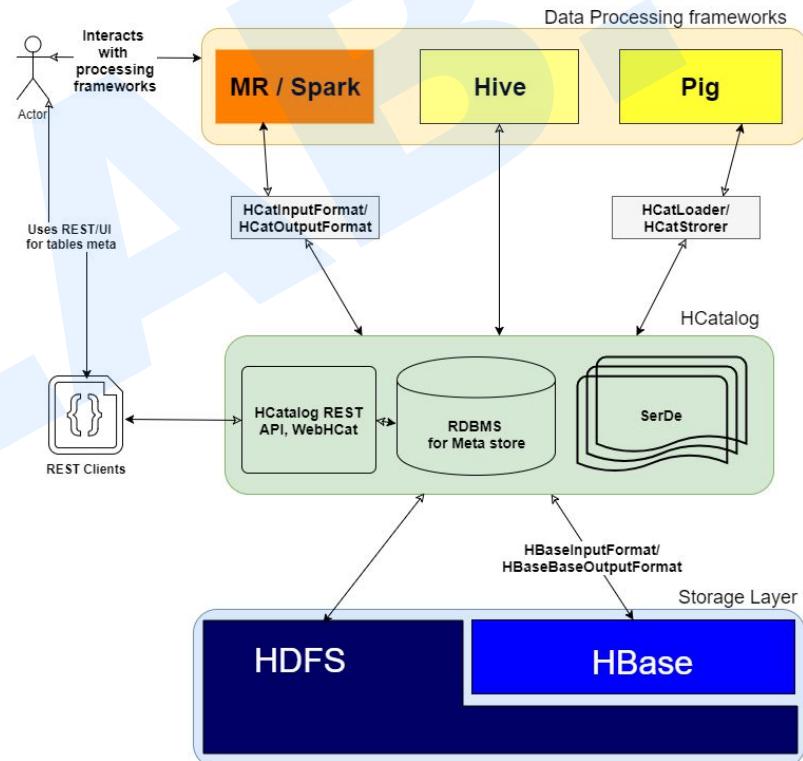
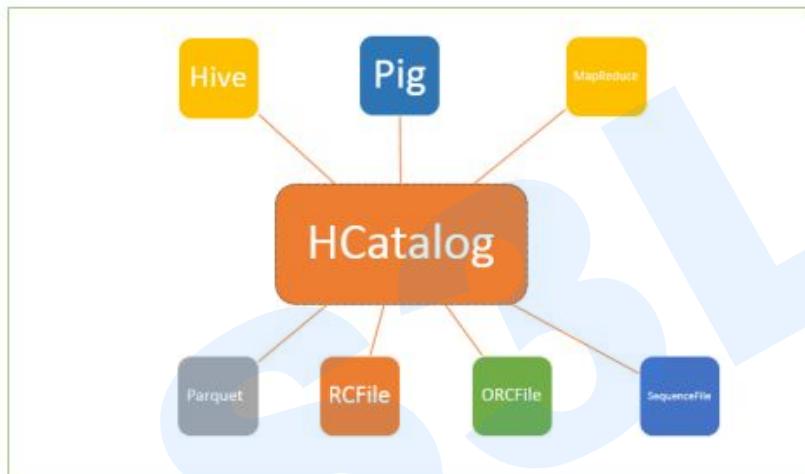
HCatalog

- A table storage management tool for Hadoop.
- It exposes the tabular data of Hive metastore to other Hadoop applications.
- It enables users with different data processing tools (Pig, MapReduce) to easily write data onto a grid. It ensures that users don't have to worry about where or in what format their data is stored.



Management Components

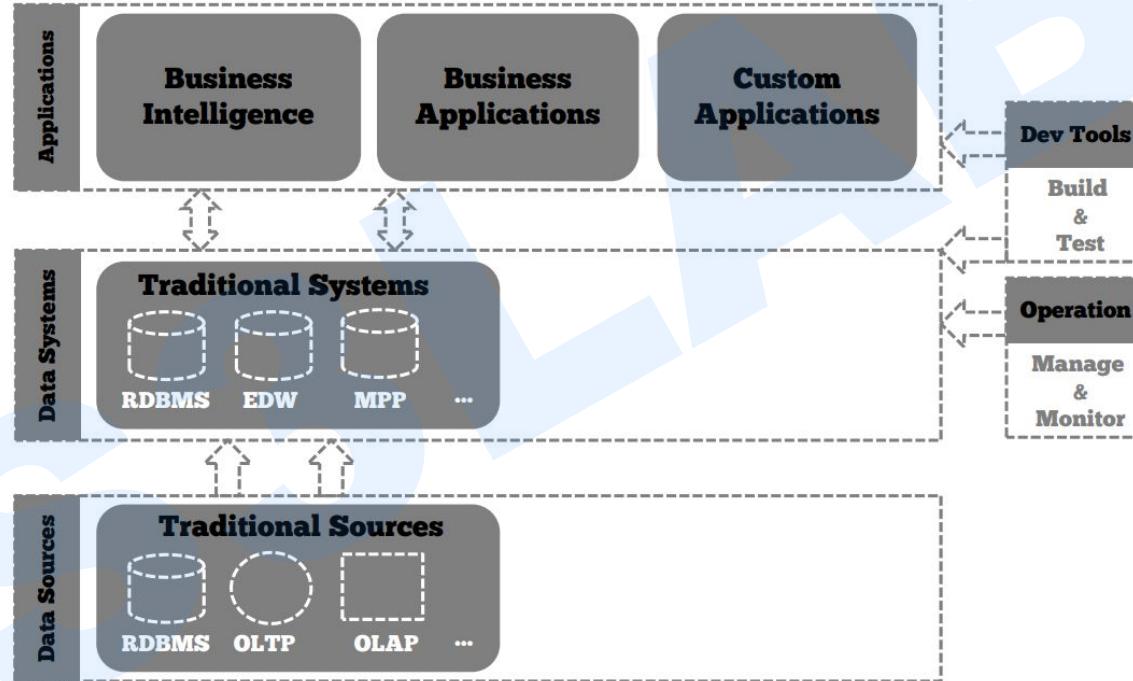
HCatalog





Use case

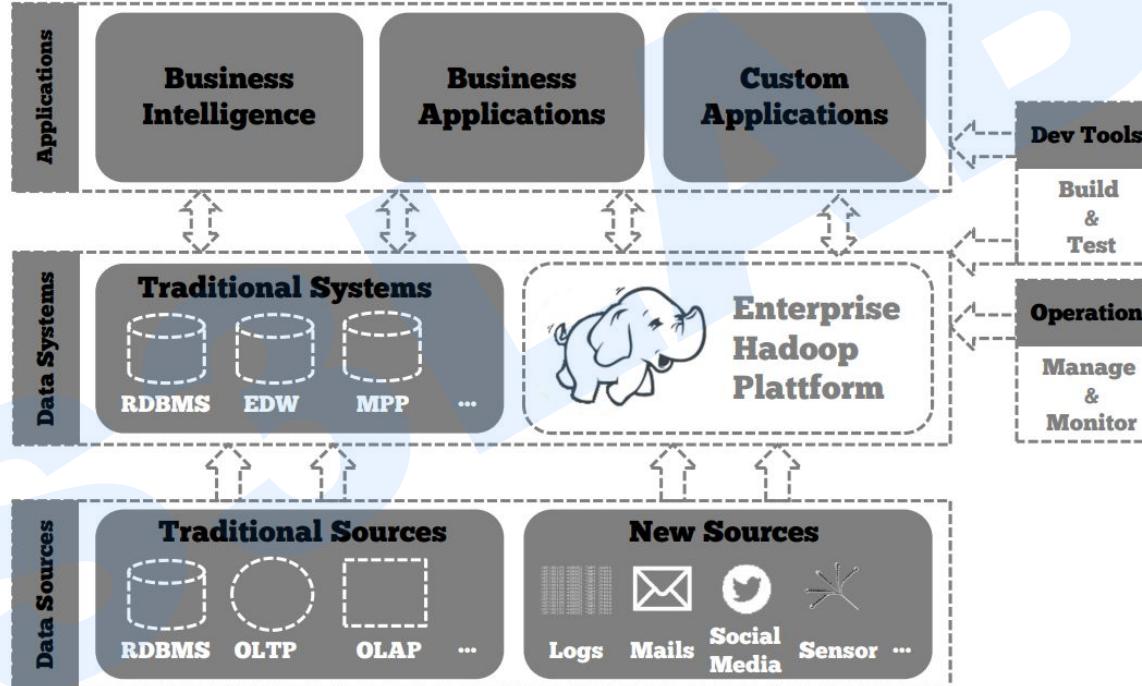
Classical enterprise platform





Use case

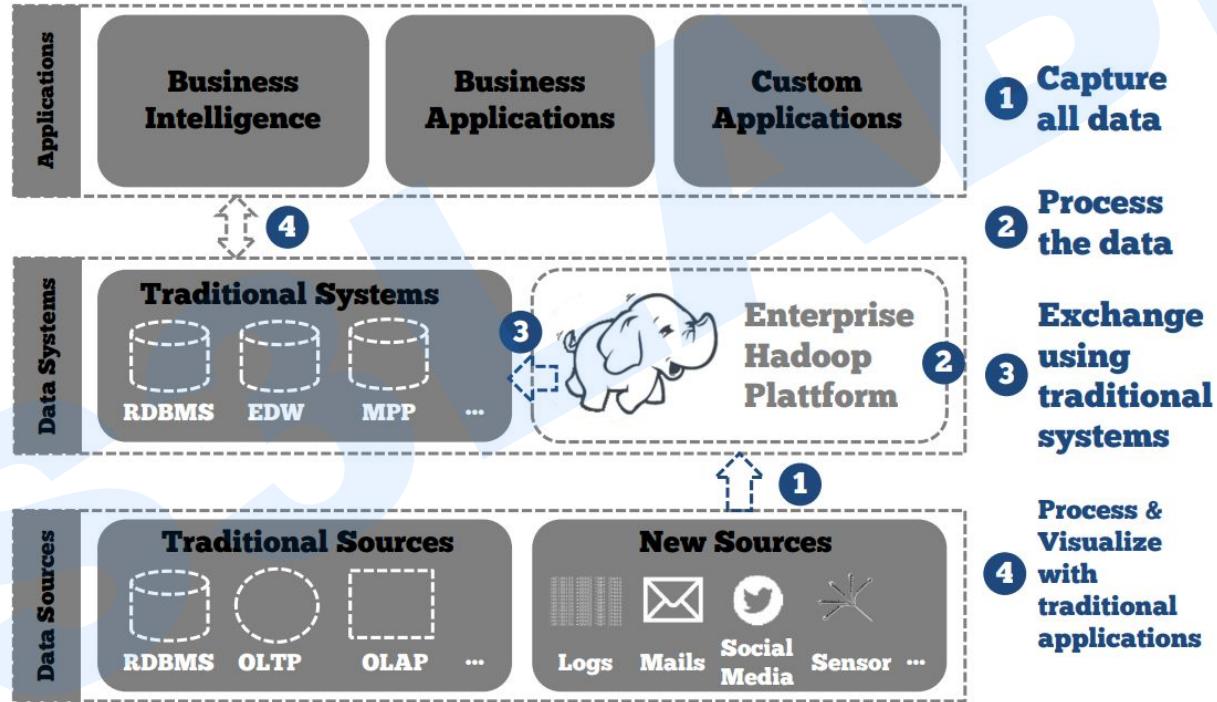
With big data





Use case

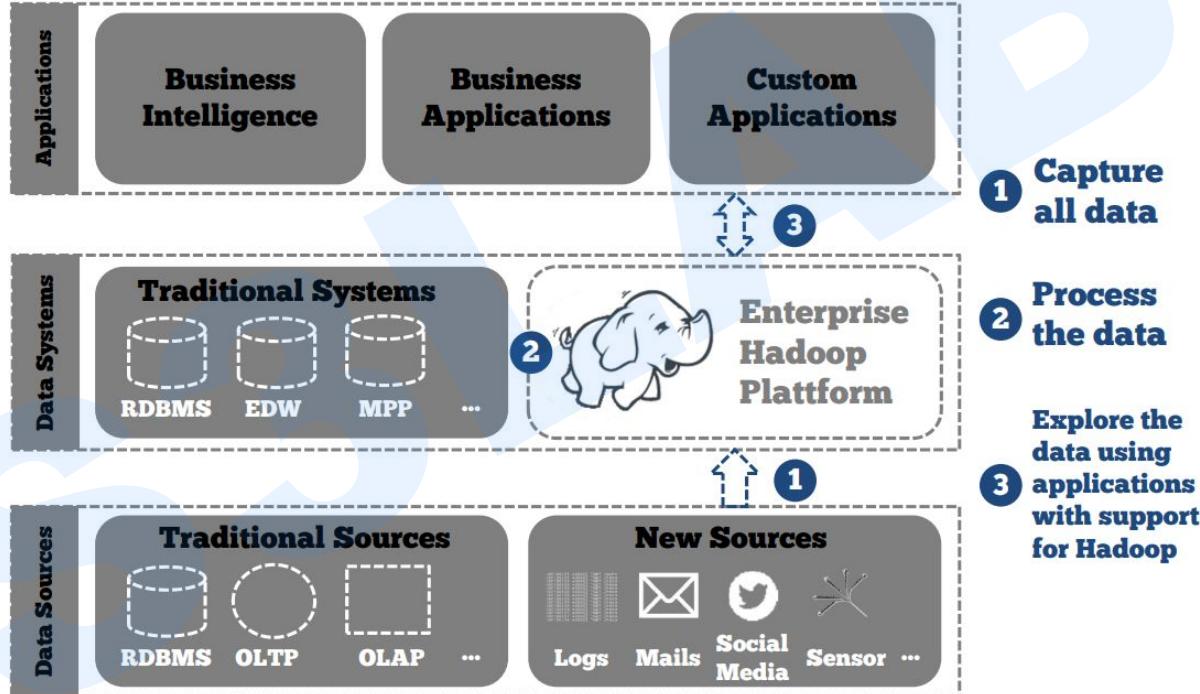
Refine data





Use case

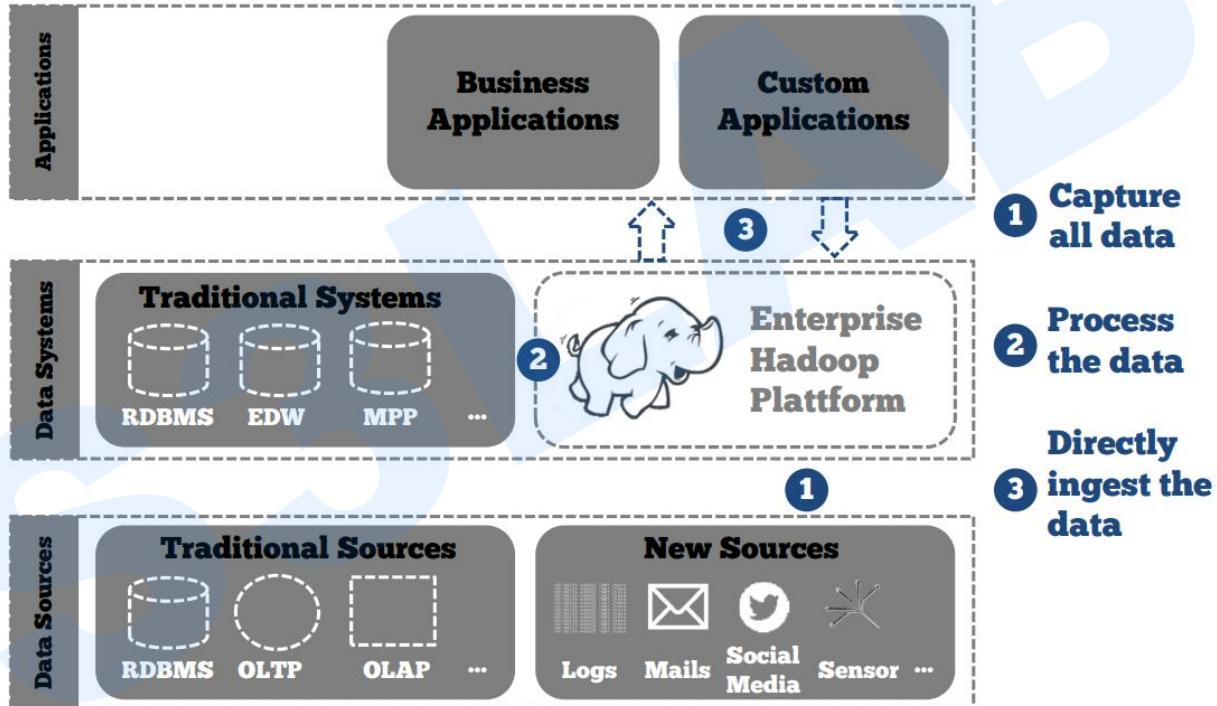
Explore data





Use case

Enrich data





An Example

Digital Advertising

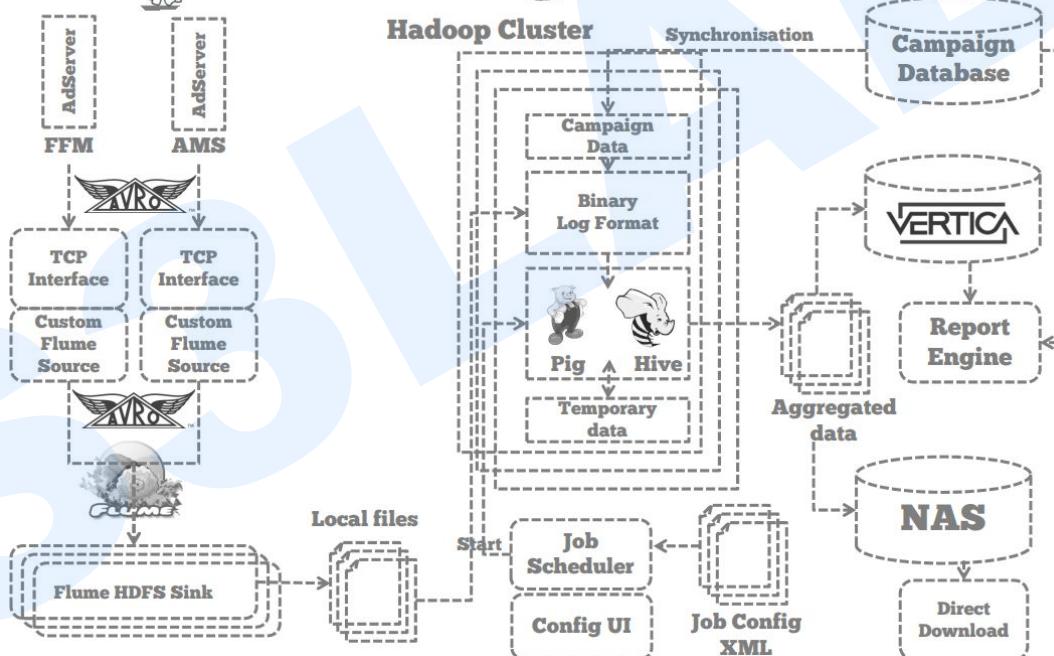
- 6 billion ad deliveries per day
- Reports (and bills) for the advertising companies needed
- Own C++ solution did not scale
- Adding functions was a nightmare



An Example

Digital Advertising

AdServing Architecture





Cảm ơn đã theo dõi

Chúng tôi hy vọng cùng nhau đi đến thành công.